

Lecture 2: Strengthening our Skills

GitHub, Python, Jetstream

Instructor: Dylan Chapp



THE UNIVERSITY OF
TENNESSEE
KNOXVILLE

Today's Topics and Problems

- Learn git and GitHub
- Build your private GitHub repository
 - Define the structure of your GitHub repository
 - Push your solution of Problem 1a in your GitHub repository
- Push the solution of assignment 1 to your private GitHub
 - If you have not yet, push the complete the solution of Assignment 1 to your private GitHub repository
- Strengthen your expertise in Python and problem solving with Assignment 2
 - Clone Assignment 2 from the course GitHub repository
 - Work on the three math problems in Assignment 2
 - Implement the solutions in Python with the Jupyter Notebook
 - Push your solution to your private GitHub repository

DUE DATE: Monday 17th before 8AM ET



Git, GitHub, and (sequential) Text Manipulation



Today we will explore:

- Learn git basics
 - Creating a git repo
 - Adding files
 - Pushing changes
 - Branching
 - Syncing & Merging
 - Undoing
- Learn GitHub basics
 - Pull Requests
- Strengthen your problem solving skills:
 - Cool math problems with different solutions, each with different performance



Git Basics

Material built from:

- <http://rogerdudler.github.io/git-guide/>
- <https://marklodato.github.io/visual-git-guide/index-en.html>



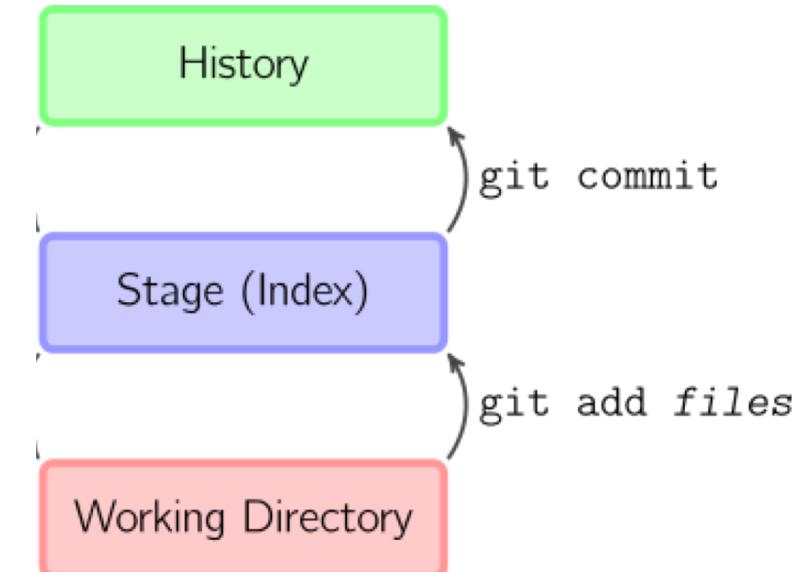
Creating a Git Repository

- If you are starting from scratch
 - `git init`
- If you already have a repo on a remote server
 - `git clone username@host:/path/to/repository`



Adding Files

- You can propose changes (add it to the **Index**) using:
 - `git add <filename>`
 - `git add .`
- You can commit these changes using:
 - `git commit -m "Commit message"`



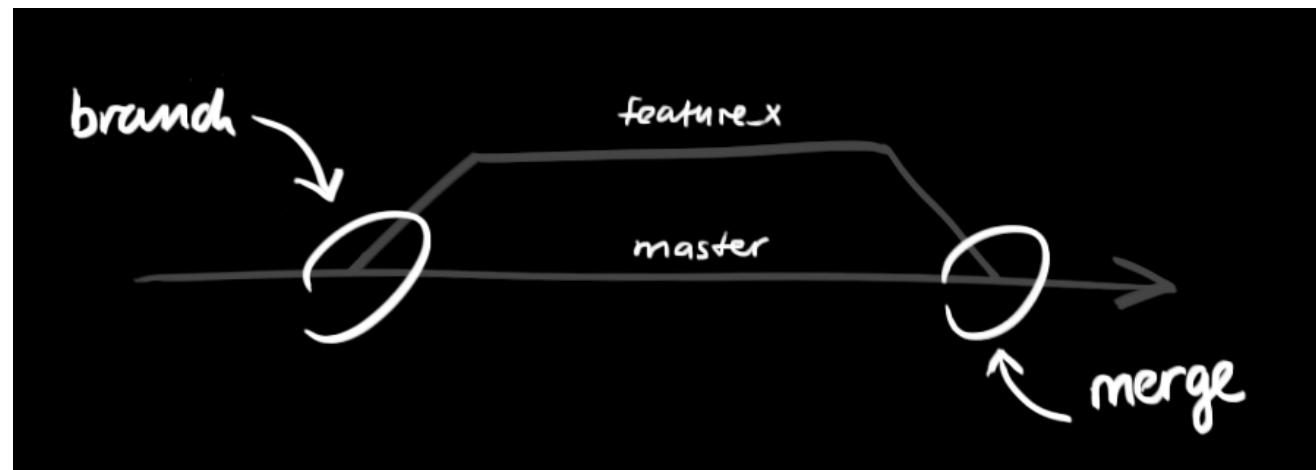
Pushing Your Changes

- Your changes are now in the **HEAD** of your local working copy
- To send those changes to your remote repository:
 - `git push origin master`
 - `git push <remote> <branch>`
- If you used ‘`git init`’, first specify a remote server:
 - `git remote add origin <server>`



Branching

- Branches are used to develop features isolated from each other.
- The *master* branch is the "default" branch when you create a repository.
- Use other branches for development and merge them back to the master branch upon completion.



Branching

- Create a new branch named "feature_x" and switch to it:
 - `git checkout -b feature_x`
- Switching back to master:
 - `git checkout master`
- Delete the feature branch:
 - `git branch -d feature_x`
- Pushing a branch for others to use:
 - `git push origin <branch>`



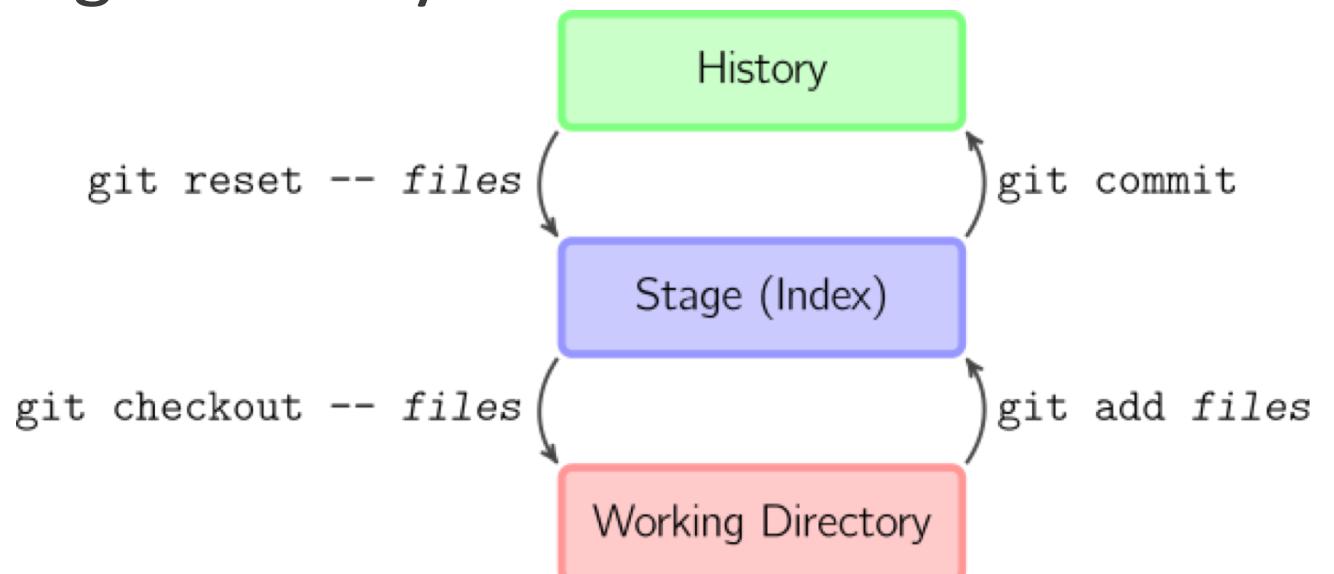
Syncing & Merging

- Updating your local repository to the newest commit:
 - `git pull`
- Merge another branch into your active branch (e.g. master):
 - `git merge <branch>`
- After resolving merge conflicts, you need to mark them as merged:
 - `git add <filename>`
- You can also preview them by using:
 - `git diff <source_branch> <target_branch>`



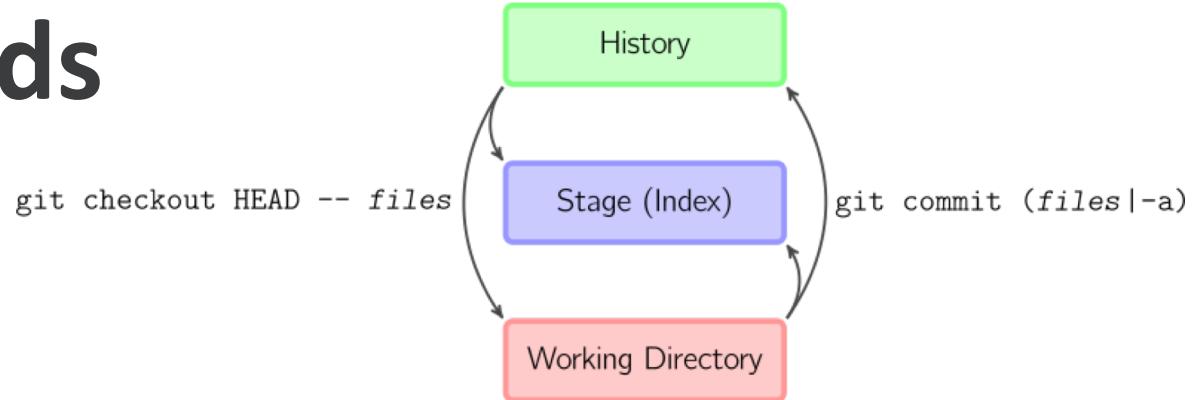
Undoing

- Unstage changes:
 - `git reset HEAD <file>...`
- Discard changes in working directory:
 - `git checkout -- <file> ...`

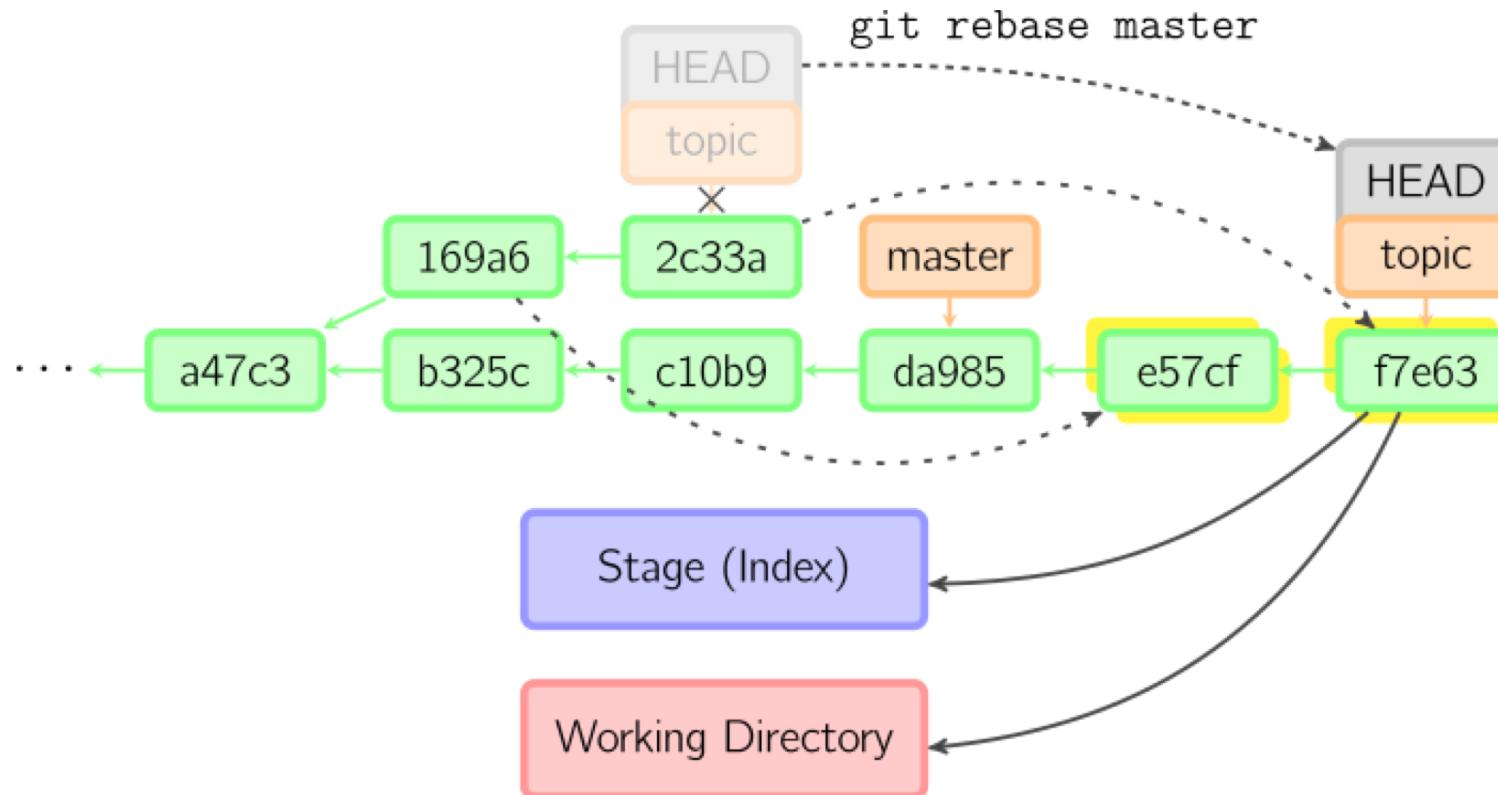


Combining Commands

- Add + Commit:
 - `git commit -a`
 - equivalent to running `git add` on all filenames that existed in the latest commit, and then running `git commit`
- Reset + Checkout:
 - `git checkout HEAD -- <file> ...`
 - copies *files* from the latest commit to both the stage and the working directory



Digging Deeper (Self-Exploration)



<https://marklodato.github.io/visual-git-guide/index-en.html>



GitHub Basics

Material built from:

<http://product.hubspot.com/blog/git-and-github-tutorial-for-beginners>



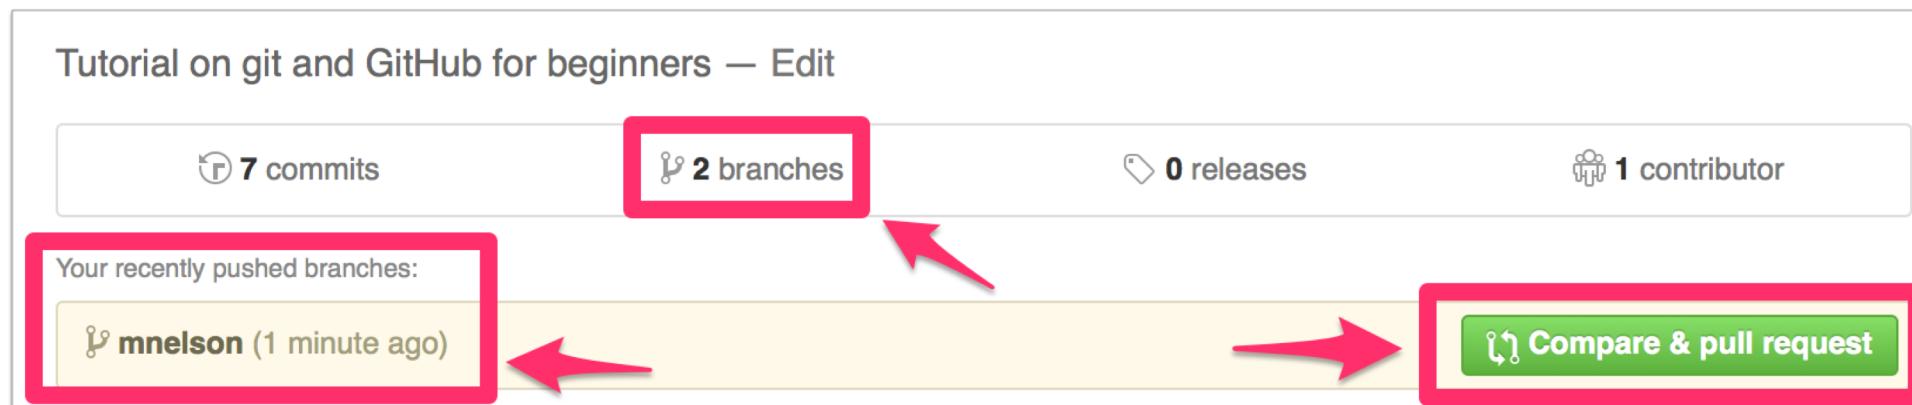
GitHub

- Cloud hosted git repositories
 - Free for open-source projects & education
 - Paid for private repositories
- Can view commits, pull requests (PRs), and issues through your browser
 - PRs are formal requests for merges of commits
 - Issues are bug reports and GitHub provides discussion threads and notifications



PRs on GitHub

- Push your code
- Visit the repo on GitHub
- Click “compare & pull request”



PRs on GitHub

Open a pull request

Create a new pull request by comparing changes across two branches. If you need to, you can also [compare across forks](#).

 base: **master** ▾ ... compare: **mnelson** ▾ ✓ Able to merge. These branches can be automatically merged.

 Added my name file

[Write](#) [Preview](#)  Markdown supported  Edit in fullscreen

Leave a comment

Attach images by dragging & dropping or [selecting them](#).



PRs on GitHub

Added my name file #1

Open cubeton wants to merge 1 commit into `master` from `mnelson`

Conversation 0 Commits 1 Files changed 1

cubeton commented a minute ago

Adding a new file with my name as the title name.

Added my name file 166b73a

Add more commits by pushing to the `mnelson` branch on `cubeton/git101`.

This branch is up-to-date with the base branch
Merging can be performed automatically.

Merge pull request

You can also [open this in GitHub for Mac](#) or view [command line instructions](#).



PRs on GitHub

Added my name file #1

Merged cubeton merged 1 commit into `master` from `mnelson` just now

Conversation 0 Commits 1 Files changed 1

 cubeton commented 7 minutes ago

Owner 

Adding a new file with my name as the title name.

 Added my name file 166b73a

 cubeton merged commit `e42a544` into `master` just now 

 Pull request successfully merged and closed 

You're all set—the `mnelson` branch can be safely deleted.



Consolidate your expertise





Assignment 2

- Consolidate expertise in Python and Jupyter Notebook
- Learn to strategize on problems' solutions
 - Write a simple solution first
 - Use git commit to save significant steps
 - Use git branch for new solution efforts



Assignment 2

Problem 1:

- Write a function that finds the sum of the multiples of p and q below n
 - A problem may have multiple solutions, some more effective than others

Problem 2:

- Manipulate a list of names
 - Use the *csv module* to import a list of names
 - Score names in a list based on name “worth” and alphabetical order



Assignment 2

Problem 3:

- Find the smallest TPH number bigger than n
 - Triangular (T), Pentagonal (P), and Hexagonal (H) numbers are generated by given formulae
 - Write the code to find the next triangle number that is also pentagonal and hexagonal







a high-level introduction

NSF Funding Areas in HPC

Traditionally concentrated on enabling petascale capability

- Blue Waters – 13.3 petaflops, 2012
- Stampede – 9.6 petaflops, 2013
- Comet – ~2.0 petaflops, 2014

Has funded research into building clouds and computer science

- CloudLab
- Chameleon

Now funding clouds to do research

- Bridges (Hybrid system)
- Jetstream



From slides by Jeremy Fischer

jetstream-cloud.org/files/Jetstream-SIAM-CSE-Presentation.pdf

What is Jetstream and why does it exist?

- NSF's first production cloud facility
- Part of the NSF eXtreme Digital (XD) program
- Provides on-demand *interactive* computing and analysis
- Enables *configurable* environments and architectures
- User-friendly, widely accessible cloud environment
- User-selectable library of preconfigured virtual machines



From slides by Jeremy Fischer
jetstream-cloud.org/files/Jetstream-SIAM-CSE-Presentation.pdf

Cloud vs HPC

Adapting to a different environment:

- No reservations, no queueing
- More interactive use and less/no batch queuing
- What? No parallel filesystem?!?
- Being your own admin – hey, we have root!
- You really can have almost any (linux) software you want**
- Constantly getting new features (<https://www.openstack.org/software/project-navigator/>)

** Here there be dragons...



From slides by Jeremy Fischer
jetstream-cloud.org/files/Jetstream-SIAM-CSE-Presentation.pdf

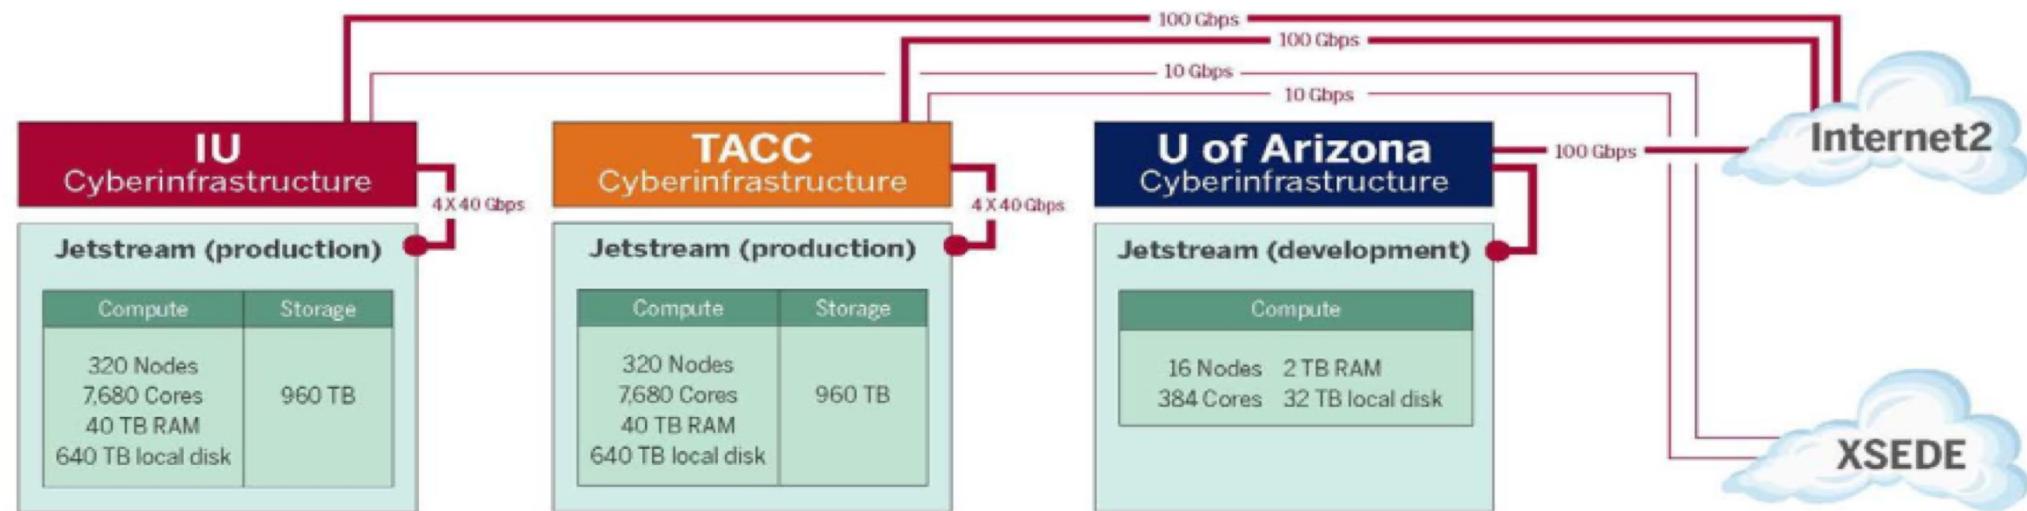
Who uses Jetstream?

- The researcher needing a handful of cores (1 to 44/vCPU)
- Software creators and researchers needing to create their own customized virtual machines and workflows
- Science gateway creators using Jetstream as either the frontend or processor for scientific jobs
- STEM Educators teaching on a variety of subjects



From slides by Jeremy Fischer
jetstream-cloud.org/files/Jetstream-SIAM-CSE-Presentation.pdf

Jetstream System Overview



From slides by Jeremy Fischer
jetstream-cloud.org/files/Jetstream-SIAM-CSE-Presentation.pdf

Hardware and Instance "Flavors"

VM Host Configuration

- Dual Intel E-2680v3 "Haswell"
- 24 physical cores/node @ 2.5 GHz (Hyperthreading on)
- 128 GB RAM
- Dual 1 TB local disks
- 10GB dual uplink NIC
- Running KVM Hypervisor

Flavor	vCPUs	RAM	Storage	Per Node
m.tiny	1	2	8	46
m.small	2	4	20	23
m.medium	6	16	60	7
m.large	10	30	120	4
m.xlarge	24	60	240	2
m.xxlarge	44	120	480	1

- Short-term storage comes as part of launched instance
- Long-term storage is XSEDE-allocated
- Implemented on backend as OpenStack Volumes
- Each user gets 10 volumes up to 500GB total storage
- Piloting object storage as well after recent update



From slides by Jeremy Fischer
jetstream-cloud.org/files/Jetstream-SIAM-CSE-Presentation.pdf

Requesting access to Jetstream

- You can request startup allocations anytime. (Startups are simple!)
- You can request allocations for educational use anytime.



From slides by Jeremy Fischer

jetstream-cloud.org/files/Jetstream-SIAM-CSE-Presentation.pdf

- We'll send instructions for making an XCede account the week before we start using Jetstream in class.
- For now, back to Assignment 02!