

Lecture 3:

The MapReduce

Programming Model

Instructor: Michela Taufer



THE UNIVERSITY OF
TENNESSEE
KNOXVILLE

BIG ORANGE. BIG IDEAS.®

Today Outline

- Learn a new programming models: MapReduce
 - MapReduce programming model: principles and definitions
 - Word Count: a concrete example of MapReduce
 - MapReduce algorithms: different approaches to solve the same problem
 - Data: Where is the data kept?
- Continue building our expertise with Jupyter and Python
 - Sequential manipulation of a classical in literature
 - Visualization of statistics

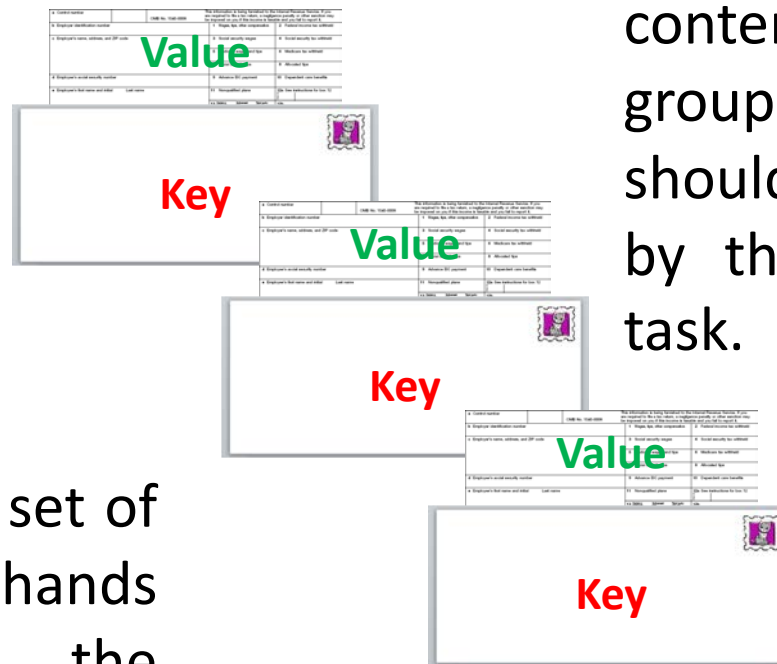
The Canonical MapReduce Programing Model

The MapReduce Workflow



Step 1: Map

The **map** task creates a set of **Key-Value** pairs and hands them off to the **sorting/shuffling** task.



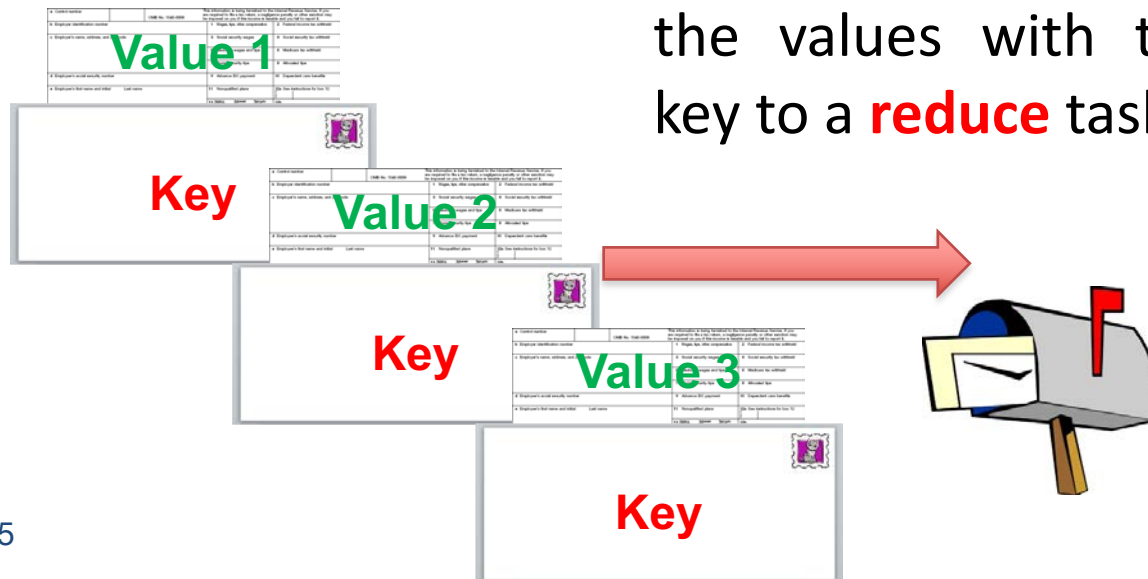
The **Value** is the content, the **Key** groups content that should be processed by the same **reduce** task.

The MapReduce Workflow



Step 2: Sort/Shuffle

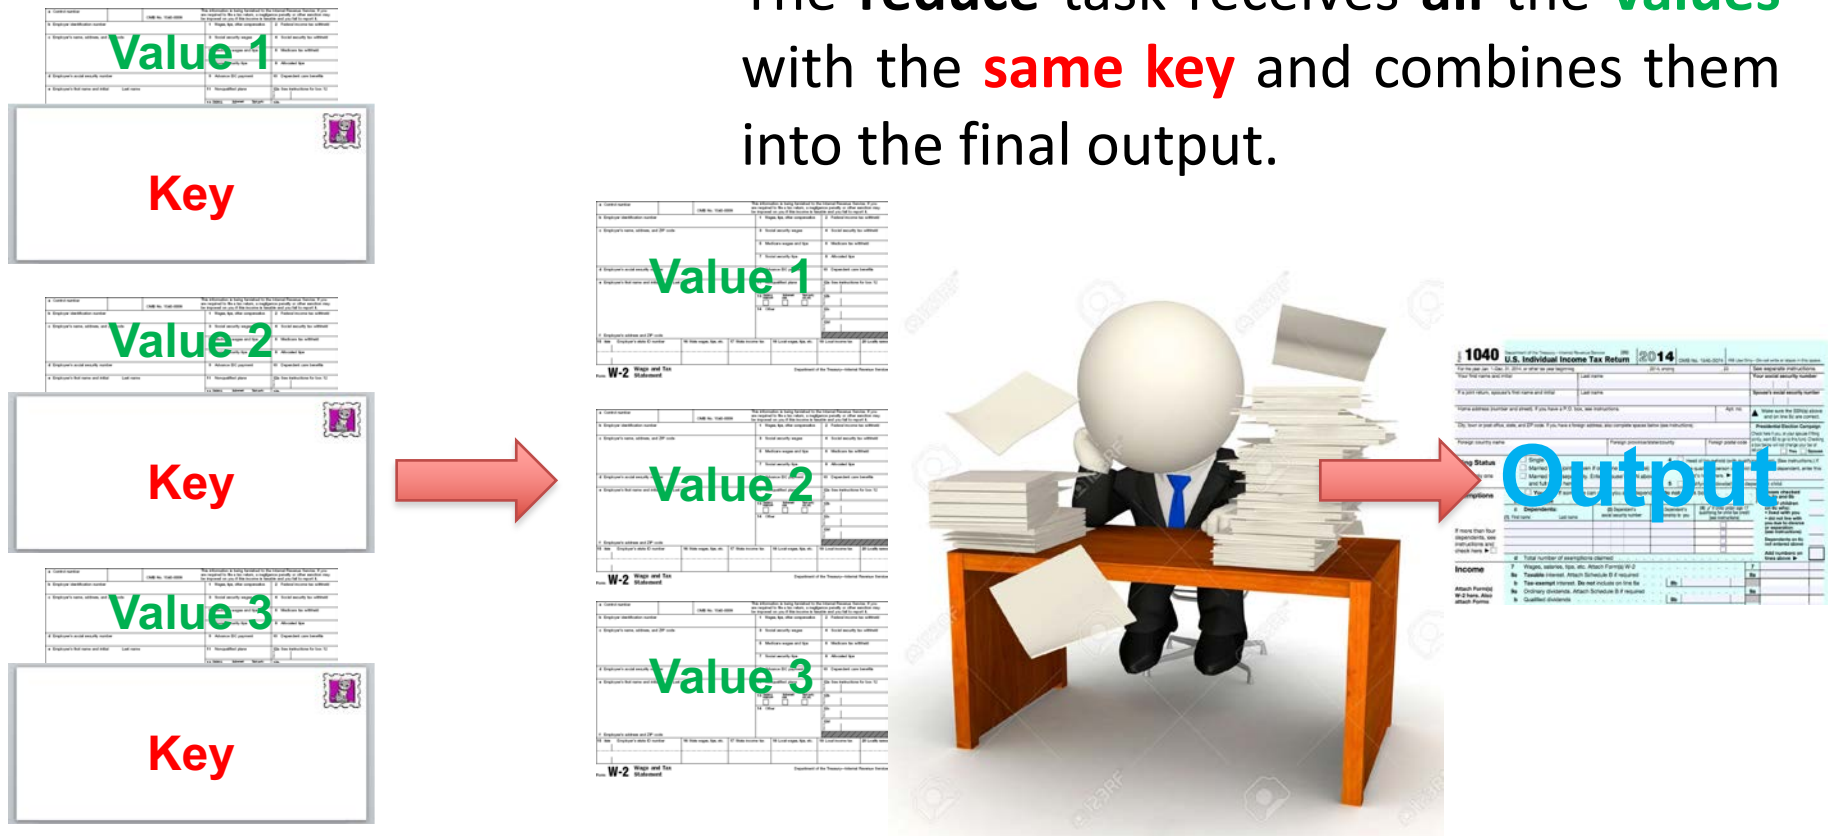
The **sorting** and **shuffling** tasks collect all the **key-value** pairs and deliver all the values with the same key to a **reduce** task.




The MapReduce Workflow

Step 3: Reduce

The **reduce** task receives **all** the **values** with the **same key** and combines them into the final output.



The MapReduce Workflow



Map

- The map function is implemented by the user.
- Each map task processes a single line of an input file at a time.
- Map tasks do not communicate with other map tasks.
- Map tasks communicate with reduce tasks **only** through the content of the value in the KV pair.
- A single map task may emit **any number** of KV pairs (including none).


Sort/Shuffle

- The sorting and shuffling process is implemented at the framework level and is opaque to the user.

Reduce

- The reduce function is implemented by the user.
- Each reduce task receives **all** the **values** associated with a given **key**.
- Reduce tasks do not communicate with each other
- A single reduce task may emit **any number** of result values for each **key** it processes.

WordCount: an example of MapReduce



Map

Key

Value

When tweetle beetles fight, → (When, 1), (**tweetle**, **1**), (beetles, 1), (fight, 1)

it's called a tweetle beetle battle. → (it's, 1), (called, 1), (a, 1), (tweetle, 1), (beetle, 1), (battle, 1)

And when they battle in a puddle, → (And, 1), (when, 1), (they, 1), (battle, 1), (in, 1), (a, 1), (puddle, 1)

it's a tweetle beetle puddle battle. → (it's, 1), (a, 1), (tweetle, 1), (beetle, 1), (puddle, 1), (battle, 1)

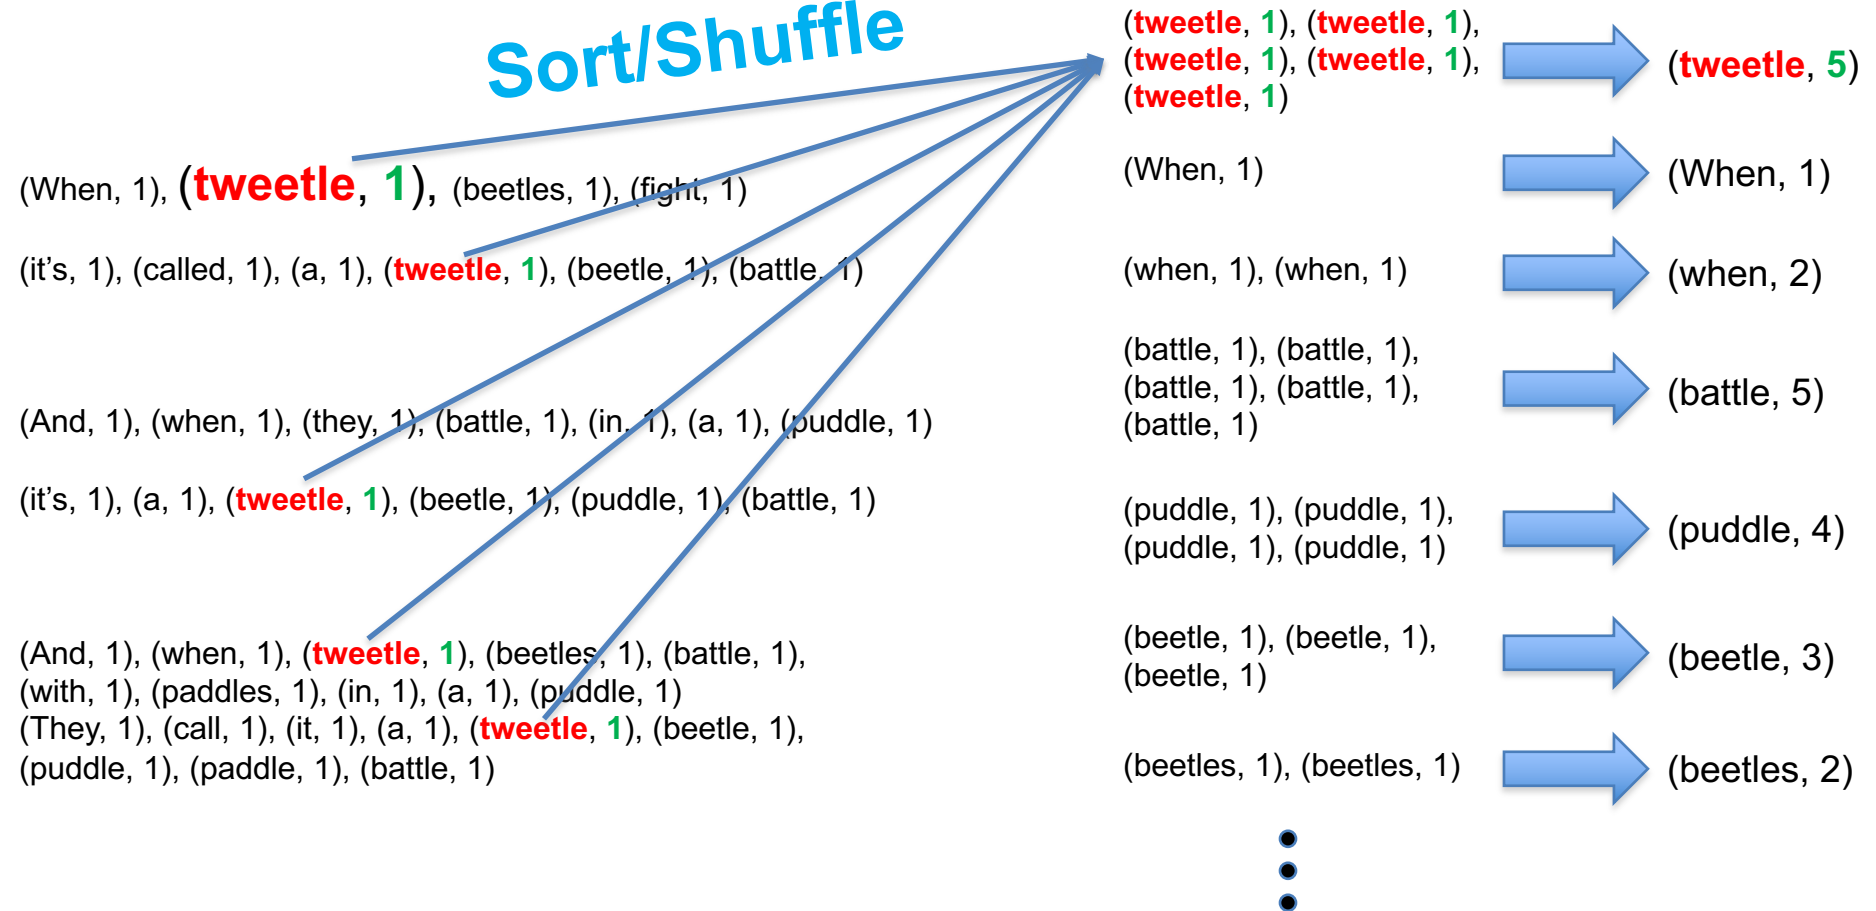
And when tweetle beetles battle with paddles in a puddle, → (And, 1), (when, 1), (tweetle, 1), (beetles, 1), (battle, 1), (with, 1), (paddles, 1), (in, 1), (a, 1), (puddle, 1)

They call it a tweetle beetle puddle paddle battle. → (They, 1), (call, 1), (it, 1), (a, 1), (tweetle, 1), (beetle, 1), (puddle, 1), (paddle, 1), (battle, 1)


WordCount: an example of MapReduce

Sort/Shuffle

Reduce




The Canonical MapReduce



Mappers: applied to all
input
data










Arbitrary number of
key-value pairs



Barrier:
distributed
sort and group by key

Shuffle and sort: aggregate values by key



Reducers: applied to
all values associated
with the same key

Pseudo-code: WordCount in MapReduce

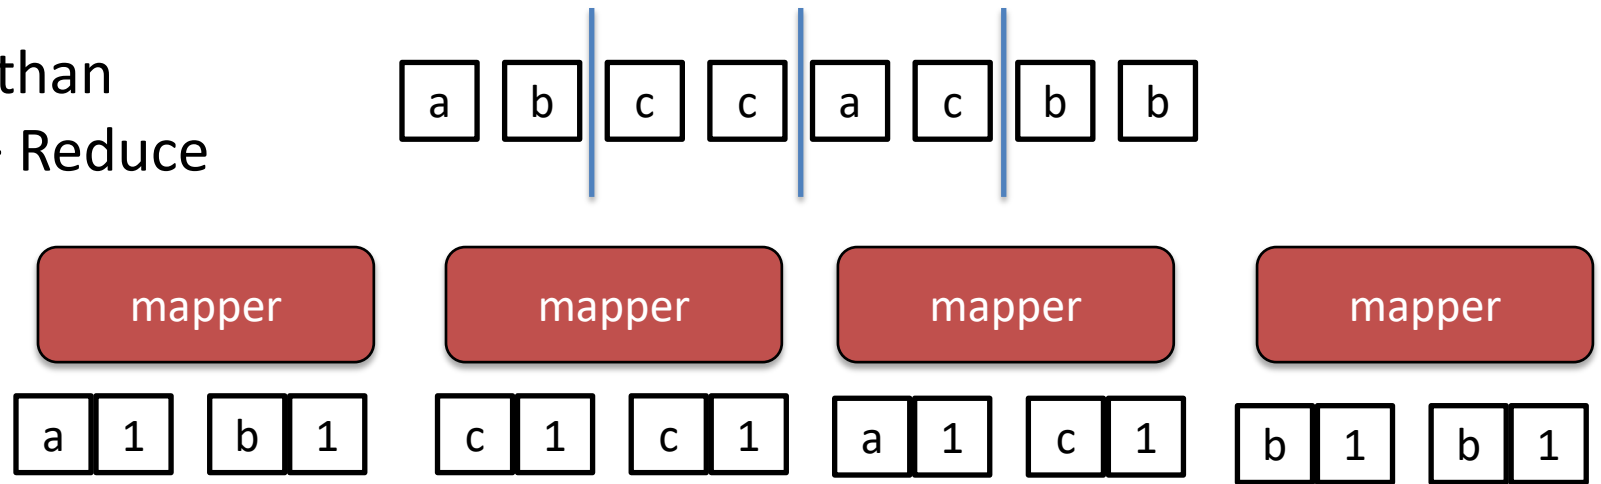
```
1: class MAPPER
2:   method MAP(docid  $a$ , doc  $d$ )
3:     for all term  $t \in \text{doc } d$  do
4:       EMIT(term  $t$ , count 1)

1: class REDUCER
2:   method REDUCE(term  $t$ , counts  $[c_1, c_2, \dots]$ )
3:      $sum \leftarrow 0$ 
4:     for all count  $c \in \text{counts } [c_1, c_2, \dots]$  do
5:        $sum \leftarrow sum + c$ 
6:     EMIT(term  $t$ , count  $sum$ )
```

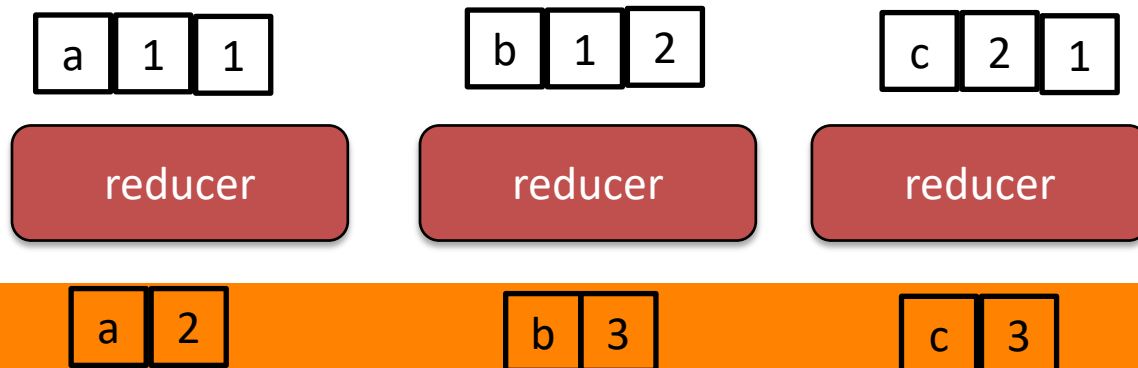
More than “Map + Reduce”

- Canonical MapReduce processing workflow:
 - Map + Reduce
- Variations:
 - Map + Combiners + Partitioners + Reduce (in Apache Hadoop and Mimir implementations)
 - Automatic implementation of combiners on the map side (in Apache Spark - more in the next lecture)

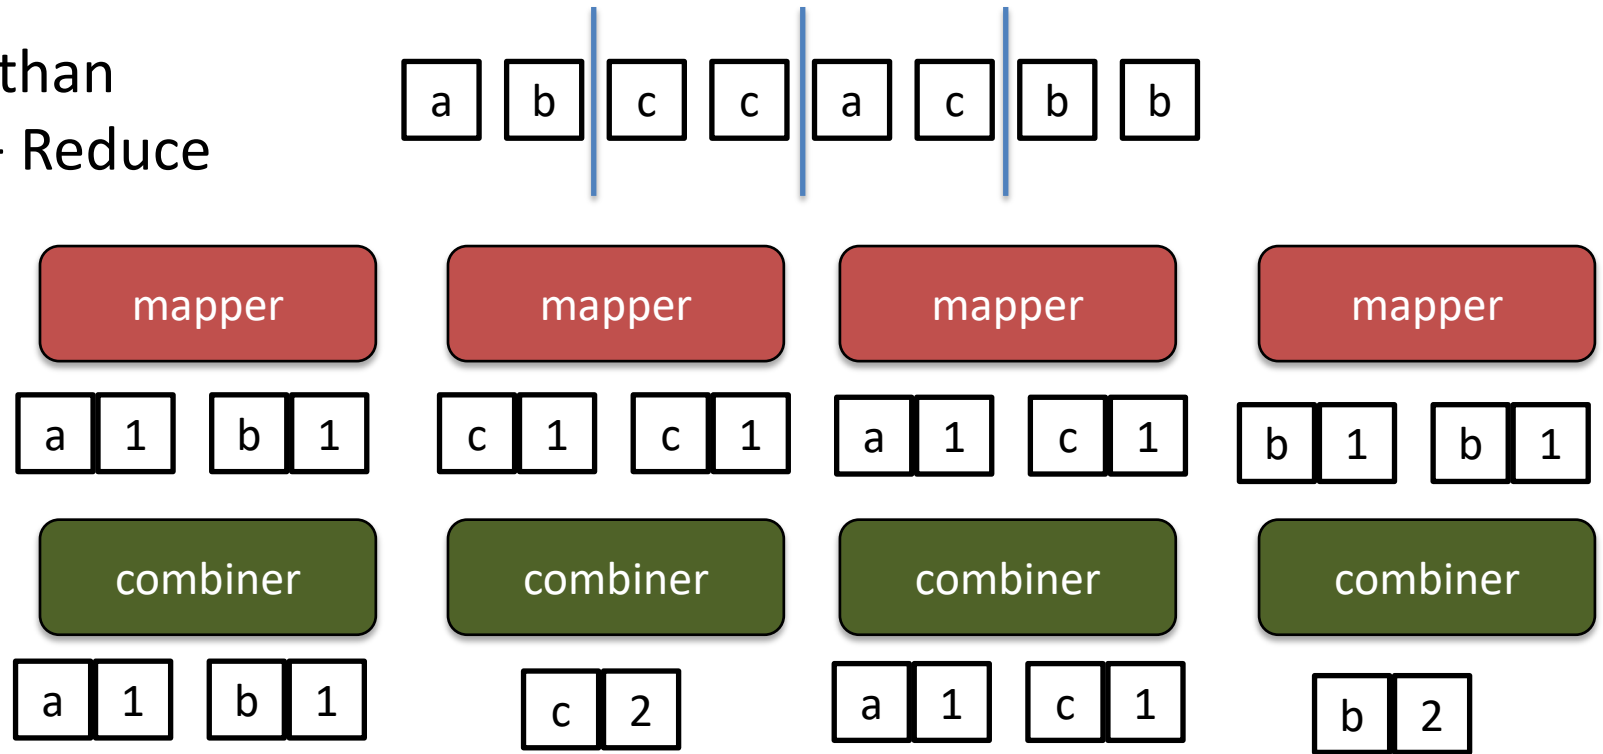
More than Map + Reduce



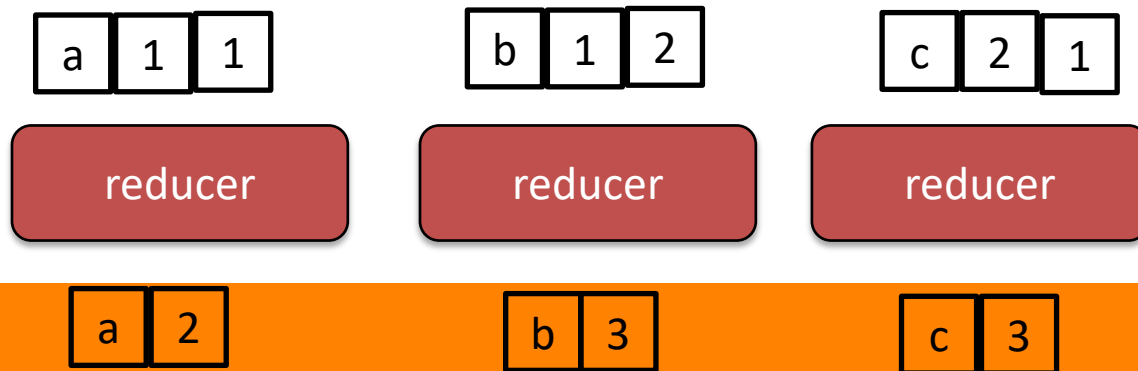
Shuffle and sort: aggregate values by key



More than Map + Reduce



Shuffle and sort: aggregate values by key

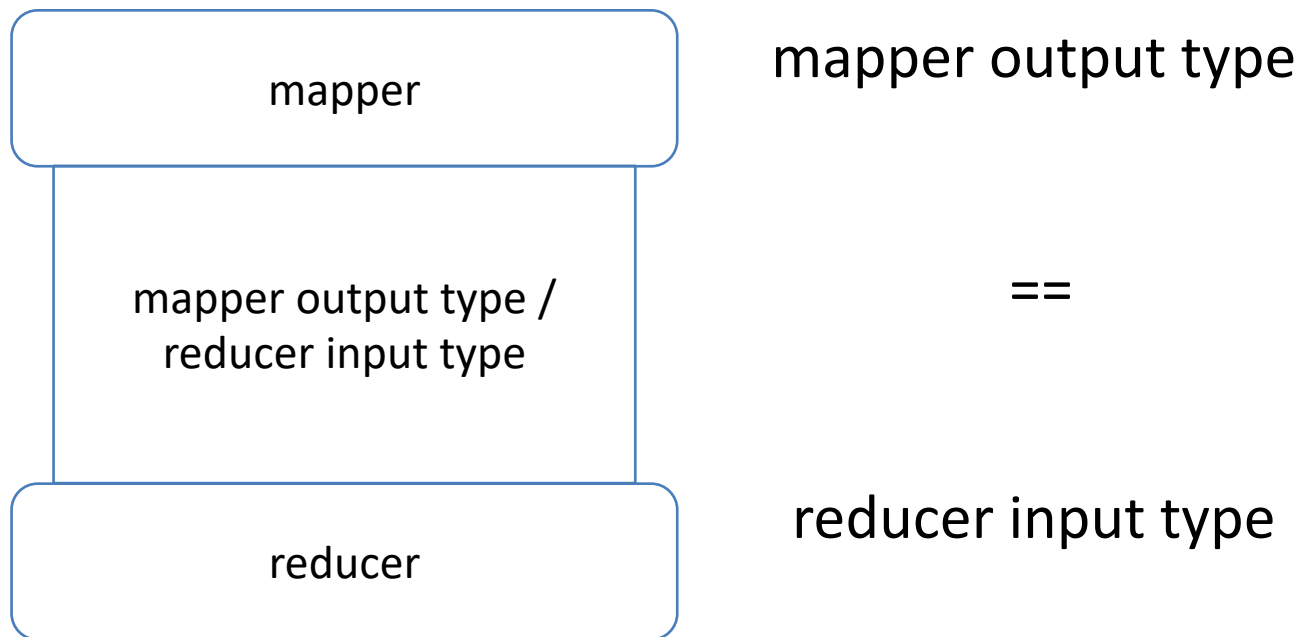


More than M + R: Combiners

- Combiners allow for local aggregation before the shuffle and sort phase
 - “mini-reducers” that take place locally, on the output of the mappers
- Strengths:
 - Number of intermediate key-value pairs moved among reducers to be **at most the number of unique words** in the collection times the number of mappers
- Weaknesses:
 - Reducers and combiners are not interchangeable unless the operation is both associative and commutative
 - Operation performed in isolation and therefore does not have access to intermediate output from other mappers

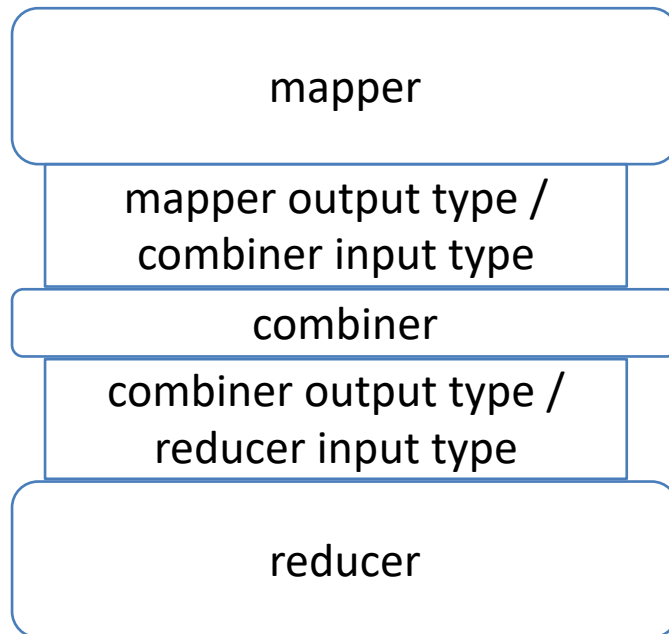
More than M + R: Combiners Constraints

- Combiners in, for example, Apache Hadoop cannot change the correctness of the MapReduce algorithm
- Combiners must have same input and output key-value types



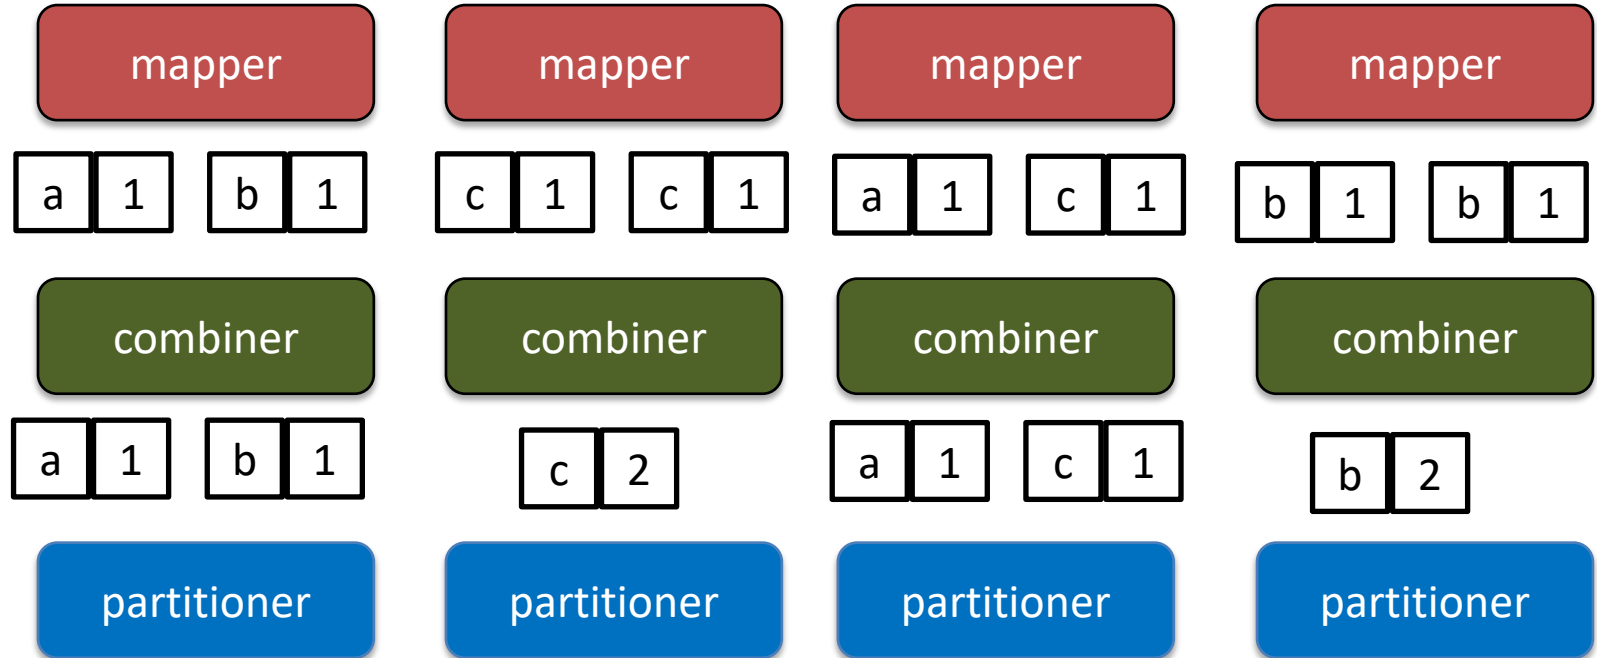

More than M + R: Combiners Constraints

- Combiners in, for example, Apache Hadoop cannot change the correctness of the MapReduce algorithm
- Combiners must have same input and output key-value types

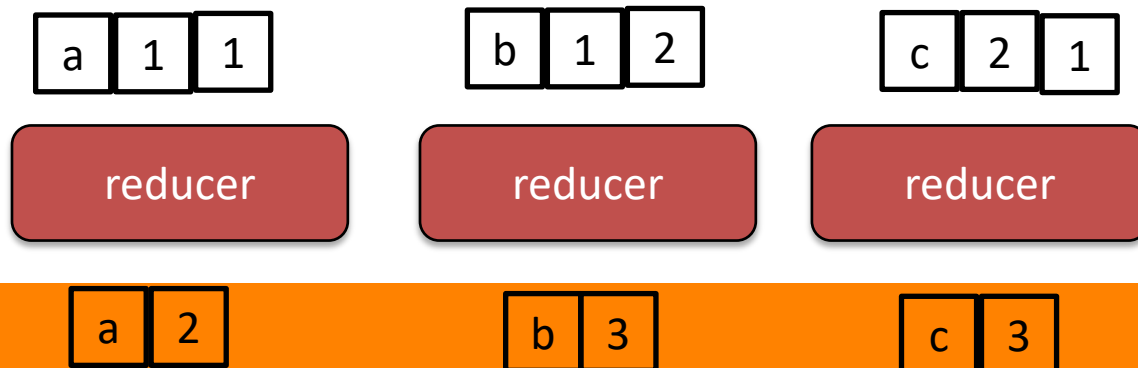


mapper output type
==
combiner input type
==
combiner output type
==
reducer input type

More than Map + Reduce



Shuffle and sort: aggregate values by key

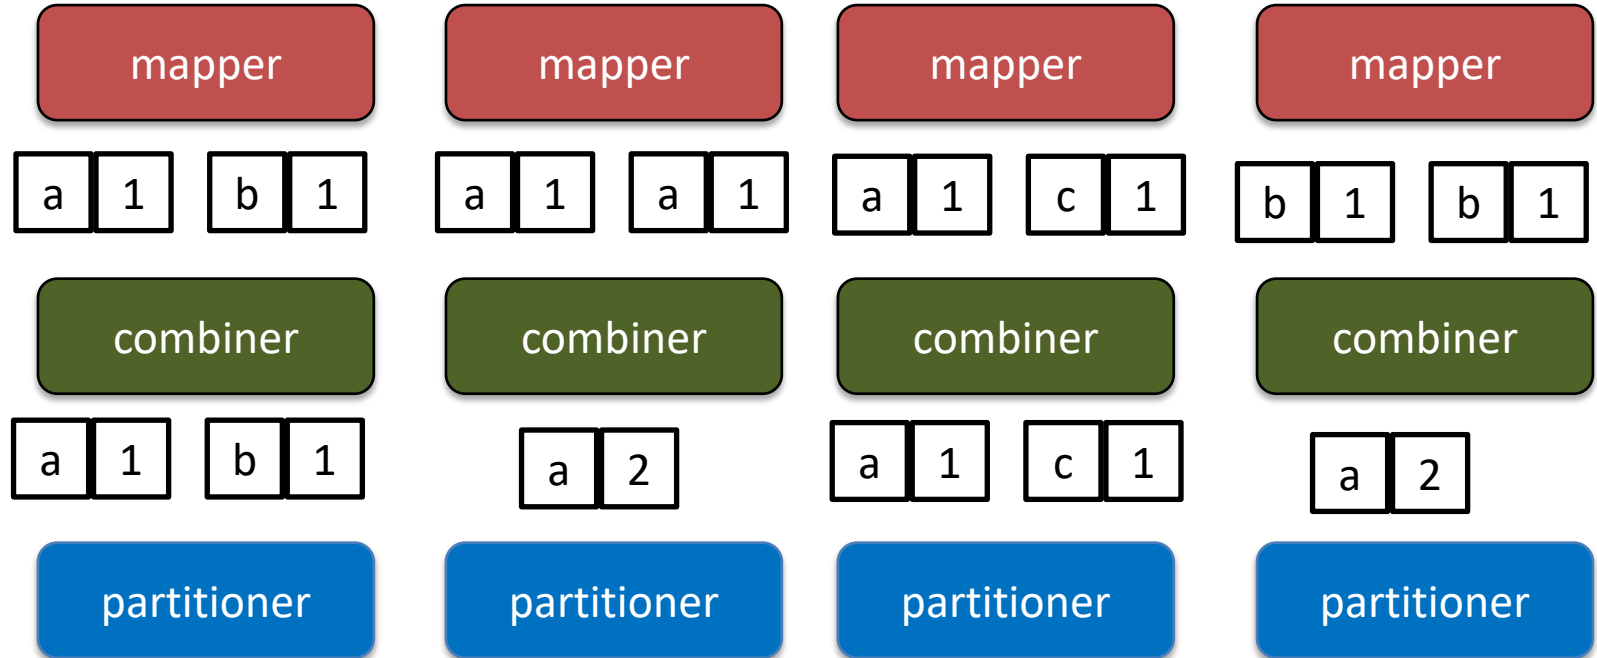


More than $M + R$: Partitioners

- Divide up the intermediate key space and **assign intermediate key-value pairs to reducers**
- Example of simple partitioner:
 - Compute hash value of a key
 - Take the mod of the value with the number of reducers
- Strengths:
 - Assign approximately the same number of keys to each reducer (dependent on the quality of the hash function)
- Weaknesses:
 - Ignore the value and different keys may have different numbers of associated values causing imbalance in the amount of data associated with each key

More than Map + Reduce

a b | c c | a c | b b



Shuffle and sort: aggregate values by key

a 1 2 1 2

b 1

c 1

reducer

reducer

reducer

a 6

b 1

c 1

WordCount in MapReduce (Review)

```
1: class MAPPER
2:     method MAP(docid  $a$ , doc  $d$ )
3:         for all term  $t \in \text{doc } d$  do
4:             EMIT(term  $t$ , count 1)

1: class REDUCER
2:     method REDUCE(term  $t$ , counts  $[c_1, c_2, \dots]$ )
3:          $sum \leftarrow 0$ 
4:         for all count  $c \in \text{counts } [c_1, c_2, \dots]$  do
5:              $sum \leftarrow sum + c$ 
6:             EMIT(term  $t$ , count  $sum$ )
```

Using Associative Array

Associative array to
aggregate term counts
on a per-document basis

```
1: class MAPPER
2:   method MAP(docid  $a$ , doc  $d$ )
3:     for all term  $t \in \text{doc } d$  do
4:       EMIT(term  $t$ , count 1)
```

```
1: class MAPPER
2:   method MAP(docid  $a$ , doc  $d$ )
3:      $H \leftarrow \text{new ASSOCIATIVEARRAY}$ 
4:     for all term  $t \in \text{doc } d$  do
5:        $H\{t\} \leftarrow H\{t\} + 1$ 
6:     for all term  $t \in H$  do
7:       EMIT(term  $t$ , count  $H\{t\}$ )
```

In-mapper Combining

```
1: class MAPPER
2:   method MAP(docid  $a$ , doc  $d$ )
3:      $H \leftarrow$  new ASSOCIATIVEARRAY
4:     for all term  $t \in$  doc  $d$  do
5:        $H\{t\} \leftarrow H\{t\} + 1$ 
6:     for all term  $t \in H$  do
7:       EMIT(term  $t$ , count  $H\{t\}$ )
```

```
1: class MAPPER
2:   method INITIALIZE
3:      $H \leftarrow$  new ASSOCIATIVEARRAY
4:   method MAP(docid  $a$ , doc  $d$ )
5:     for all term  $t \in$  doc  $d$  do
6:        $H\{t\} \leftarrow H\{t\} + 1$ 
7:   method CLOSE
8:     for all term  $t \in H$  do
9:       EMIT(term  $t$ , count  $H\{t\}$ )
```

- Initialize an associative array for holding term counts
- Accumulate partial counts in associative array across multiple documents
- incorporate combiner functionality directly inside the mapper (in-mapper combining)

**One problems,
many solutions**

A real problem

- Given a very large log report from a soccer website with the *number of scores per game* of all the soccer players worldwide for seasons 2013 – 2018
 - Keys represent soccer player name
 - Values represent the number of scores a player get per game
 - Data is chronologically sorted based on the date of the game
- **Which players shall be awarded the Soccer's Best Player Award?**
 - Criteria: the winner has the highest mean number of scores per played game
- Compute the mean number of scores on a per-player basis

Things to remember ...

Mean(1; 2; 3; 4; 5) **IS NOT** Mean(Mean(1; 2); Mean(3; 4; 5))

A solution

```
1: class MAPPER
2:     method MAP(string  $t$ , integer  $r$ )
3:         EMIT(string  $t$ , integer  $r$ )

1: class REDUCER
2:     method REDUCE(string  $t$ , integers  $[r_1, r_2, \dots]$ )
3:          $sum \leftarrow 0$ 
4:          $cnt \leftarrow 0$ 
5:         for all integer  $r \in$  integers  $[r_1, r_2, \dots]$  do
6:              $sum \leftarrow sum + r$ 
7:              $cnt \leftarrow cnt + 1$ 
8:          $r_{avg} \leftarrow sum / cnt$ 
9:         EMIT(string  $t$ , integer  $r_{avg}$ )
```

A solution

Q: Does the solution work?

```
1: class MAPPER
2:     method MAP(string  $t$ , integer  $r$ )
3:         EMIT(string  $t$ , integer  $r$ )

1: class REDUCER
2:     method REDUCE(string  $t$ , integers  $[r_1, r_2, \dots]$ )
3:          $sum \leftarrow 0$ 
4:          $cnt \leftarrow 0$ 
5:         for all integer  $r \in$  integers  $[r_1, r_2, \dots]$  do
6:              $sum \leftarrow sum + r$ 
7:              $cnt \leftarrow cnt + 1$ 
8:          $r_{avg} \leftarrow sum / cnt$ 
9:         EMIT(string  $t$ , integer  $r_{avg}$ )
```

A solution

Q: Is this an efficient solution?

```
1: class MAPPER
2:     method MAP(string  $t$ , integer  $r$ )
3:         EMIT(string  $t$ , integer  $r$ )

1: class REDUCER
2:     method REDUCE(string  $t$ , integers  $[r_1, r_2, \dots]$ )
3:          $sum \leftarrow 0$ 
4:          $cnt \leftarrow 0$ 
5:         for all integer  $r \in$  integers  $[r_1, r_2, \dots]$  do
6:              $sum \leftarrow sum + r$ 
7:              $cnt \leftarrow cnt + 1$ 
8:          $r_{avg} \leftarrow sum / cnt$ 
9:         EMIT(string  $t$ , integer  $r_{avg}$ )
```

Another solution?

```
1: class MAPPER
2:     method MAP(string  $t$ , integer  $r$ )
3:         EMIT(string  $t$ , integer  $r$ )

1: class COMBINER
2:     method COMBINE(string  $t$ , integers  $[r_1, r_2, \dots]$ )
3:          $sum \leftarrow 0$ 
4:          $cnt \leftarrow 0$ 
5:         for all integer  $r \in$  integers  $[r_1, r_2, \dots]$  do
6:              $sum \leftarrow sum + r$ 
7:              $cnt \leftarrow cnt + 1$ 
8:         EMIT(string  $t$ , pair ( $sum$ ,  $cnt$ ))

1: class REDUCER
2:     method REDUCE(string  $t$ , pairs  $[(s_1, c_1), (s_2, c_2) \dots]$ )
3:          $sum \leftarrow 0$ 
4:          $cnt \leftarrow 0$ 
5:         for all pair  $(s, c) \in$  pairs  $[(s_1, c_1), (s_2, c_2) \dots]$  do
6:              $sum \leftarrow sum + s$ 
7:              $cnt \leftarrow cnt + c$ 
8:          $r_{avg} \leftarrow sum / cnt$ 
9:         EMIT(string  $t$ , integer  $r_{avg}$ )
```


Another **portable** solution?

```
1: class MAPPER
2:     method MAP(string  $t$ , integer  $r$ )
3:         EMIT(string  $t$ , integer  $r$ )

1: class COMBINER
2:     method COMBINE(string  $t$ , integers  $[r_1, r_2, \dots]$ )
3:          $sum \leftarrow 0$ 
4:          $cnt \leftarrow 0$ 
5:         for all integer  $r \in$  integers  $[r_1, r_2, \dots]$  do
6:              $sum \leftarrow sum + r$ 
7:              $cnt \leftarrow cnt + 1$ 
8:         EMIT(string  $t$ , pair ( $sum, cnt$ ))

1: class REDUCER
2:     method REDUCE(string  $t$ , pairs  $[(s_1, c_1), (s_2, c_2) \dots]$ )
3:          $sum \leftarrow 0$ 
4:          $cnt \leftarrow 0$ 
5:         for all pair  $(s, c) \in$  pairs  $[(s_1, c_1), (s_2, c_2) \dots]$  do
6:              $sum \leftarrow sum + s$ 
7:              $cnt \leftarrow cnt + c$ 
8:          $r_{avg} \leftarrow sum / cnt$ 
9:         EMIT(string  $t$ , integer  $r_{avg}$ )
```


Another **portable** solution?

```
1: class MAPPER
2:     method MAP(string  $t$ , integer  $r$ )
3:         EMIT(string  $t$ , integer  $r$ )
```

```
1: class COMBINER
2:     method COMBINE(string  $t$ , integers [ $r_1, r_2, \dots$ ])
3:          $sum \leftarrow 0$ 
4:          $cnt \leftarrow 0$ 
5:         for all integer  $r \in$  integers [ $r_1, r_2, \dots$ ] do
6:              $sum \leftarrow sum + r$ 
7:              $cnt \leftarrow cnt + 1$ 
8:         EMIT(string  $t$ , pair ( $sum, cnt$ ))
```

```
1: class REDUCER
2:     method REDUCE(string  $t$ , pairs [ $(s_1, c_1), (s_2, c_2) \dots$ ])
3:          $sum \leftarrow 0$ 
4:          $cnt \leftarrow 0$ 
5:         for all pair  $(s, c) \in$  pairs [ $(s_1, c_1), (s_2, c_2) \dots$ ] do
6:              $sum \leftarrow sum + s$ 
7:              $cnt \leftarrow cnt + c$ 
8:          $r_{avg} \leftarrow sum / cnt$ 
9:         EMIT(string  $t$ , integer  $r_{avg}$ )
```

Another **portable** solution?

```
1: class MAPPER
2:     method MAP(string t, integer r)
3:         EMIT(string t, integer r)

1: class COMBINER
2:     method COMBINE(string t, integers [r1, r2, ...])
3:         sum ← 0
4:         cnt ← 0
5:         for all integer r ∈ integers [r1, r2, ...] do
6:             sum ← sum + r
7:             cnt ← cnt + 1
8:         EMIT(string t, pair (sum, cnt))

1: class REDUCER
2:     method REDUCE(string t, pairs [(s1, c1), (s2, c2) ...])
```

Mismatch between combiner input key-value type and output key-value type violates the MapReduce programming model!!!

A portable and efficient solution

```
1: class MAPPER
2:     method MAP(string  $t$ , integer  $r$ )
3:         EMIT(string  $t$ , pair ( $r$ , 1))

1: class COMBINER
2:     method COMBINE(string  $t$ , pairs  $[(s_1, c_1), (s_2, c_2) \dots]$ )
3:          $sum \leftarrow 0$ 
4:          $cnt \leftarrow 0$ 
5:         for all pair  $(s, c) \in$  pairs  $[(s_1, c_1), (s_2, c_2) \dots]$  do
6:              $sum \leftarrow sum + s$ 
7:              $cnt \leftarrow cnt + c$ 
8:         EMIT(string  $t$ , pair ( $sum$ ,  $cnt$ ))

1: class REDUCER
2:     method REDUCE(string  $t$ , pairs  $[(s_1, c_1), (s_2, c_2) \dots]$ )
3:          $sum \leftarrow 0$ 
4:          $cnt \leftarrow 0$ 
5:         for all pair  $(s, c) \in$  pairs  $[(s_1, c_1), (s_2, c_2) \dots]$  do
6:              $sum \leftarrow sum + s$ 
7:              $cnt \leftarrow cnt + c$ 
8:          $r_{avg} \leftarrow sum / cnt$ 
9:         EMIT(string  $t$ , integer  $r_{avg}$ )
```

In-mapper Combining Design Pattern Solution

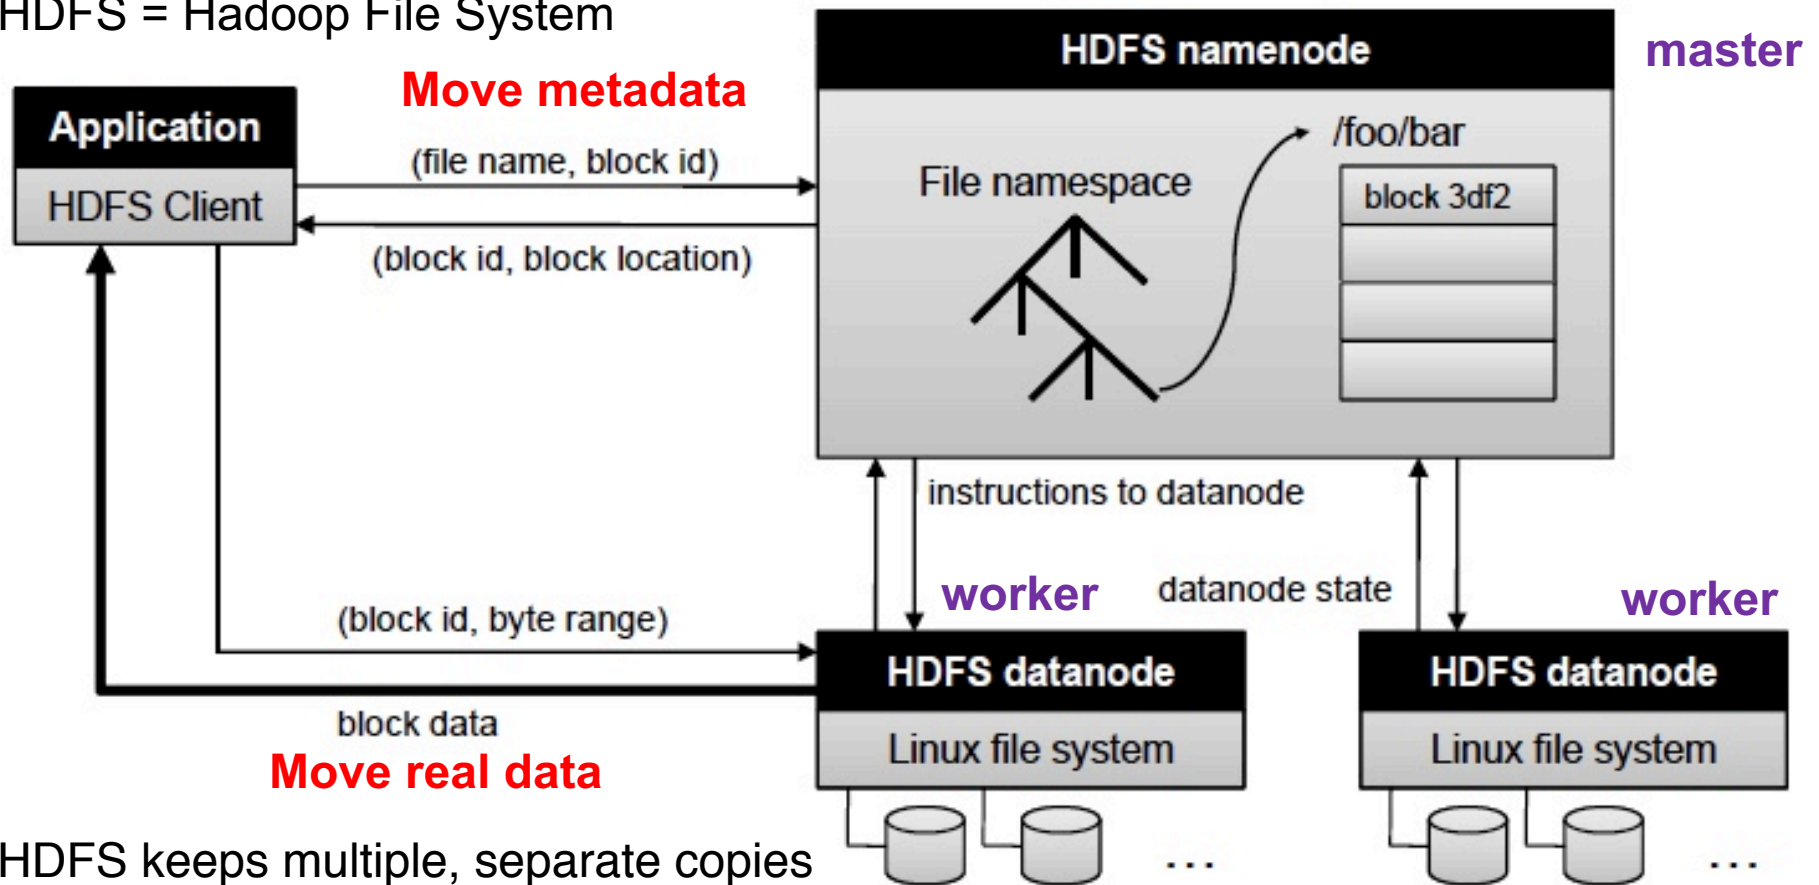
```
1: class MAPPER
2:     method INITIALIZE
3:          $S \leftarrow \text{new ASSOCIATIVEARRAY}$ 
4:          $C \leftarrow \text{new ASSOCIATIVEARRAY}$ 
5:     method MAP(string  $t$ , integer  $r$ )
6:          $S\{t\} \leftarrow S\{t\} + r$ 
7:          $C\{t\} \leftarrow C\{t\} + 1$ 
8:     method CLOSE
9:         for all term  $t \in S$  do
10:             EMIT(term  $t$ , pair ( $S\{t\}$ ,  $C\{t\}$ ))

1: class REDUCER
2:     method REDUCE(string  $t$ , pairs  $[(s_1, c_1), (s_2, c_2) \dots]$ )
3:          $sum \leftarrow 0$ 
4:          $cnt \leftarrow 0$ 
5:         for all pair  $(s, c) \in \text{pairs } [(s_1, c_1), (s_2, c_2) \dots]$  do
6:              $sum \leftarrow sum + s$ 
7:              $cnt \leftarrow cnt + c$ 
8:          $r_{avg} \leftarrow sum / cnt$ 
9:         EMIT(string  $t$ , integer  $r_{avg}$ )
```

Where is the data kept?

The Hadoop File System

HDFS = Hadoop File System



HDFS keeps multiple, separate copies of each data block to ensure reliability, availability, and performance

Actual data organized in blocks

Assignment 3

Assignment 3 - CS 594 / CS 690

- Goal: Continue building our expertise with Jupyter and Python
 - Sequential manipulation of a classical in literature
 - Visualization of statistics
- Deadline: ***September 24, 8AM ET***

Assignment 3 - CS 594 / CS 690

- Given a literature classic such as the “The Count of Monte Cristo”
- Problem 1: Analyze the text for word length frequency
- Problem 2: Analyze the text for letter frequency
- Problem 3: Count the positional frequencies of each letter (first, interior, and last)
- Problem 4: Visualize your findings in histograms (one for each one of Problems 1-3)
 - One code is give to you, write the other two codes

Deadline: September 24 - 8AM ET

Assignment 3 - CS 690

- Read paper “MapReduce: Simplified Data Processing on Large Clusters” *by* Jeffrey Dean and Sanjay Ghemawat, Google Inc.
- Submit summary:
 - Add summary to your private GitHub repository
 - Use the template provided
 - Follow mandatory requirements for your summary

Deadline: October 1 - 8AM ET



THE UNIVERSITY OF
TENNESSEE
KNOXVILLE

BIG ORANGE. BIG IDEAS.®