

Lecture 4:

Introduction to Spark

Instructor: Michela Taufer



THE UNIVERSITY OF
TENNESSEE
KNOXVILLE


BIG ORANGE. BIG IDEAS.®

Today Outline

- Overview of Spark core concepts
- Create SparkContext and Resilient distributed datasets (RDDs)
- Run parallel operations on RDDs
- Assignment 4: create sequential Spark call:
 - *mapSequential*,
 - *reduceSequential*
 - *reduceByKeySequential*
- Work on Assignment 4
- Access our XSEDE Jetstream account
- Work on Assignment 4

Spark reference

O'REILLY®



Holden Karau,
Andy Kowinski & Matei Zaharia

Spark is ...

- ... a “computational engine that is responsible for scheduling, distributing, and monitoring applications consisting of many computational tasks across many **worker** machines, or a computing cluster.”

Running Spark

- Apache Spark has four APIs: Java, Scala, R, and **Python**
- Ways to use the Python API:
 - Interactively, using **pyspark**
 - Non-interactively, using **spark-submit**
 - Within the **Jupyter Notebook**
- Run interactively: open the Python version of the Spark shell
bin/pyspark
- Run non-interactively: Write Python scripts and run script:
bin/spark-submit my_script.py
- Run in Jupyter Notebook: Assignment 5

Running Spark

- Apache Spark has four APIs: Java, Scala, R, and **Python**
- Ways to use the Python API:
 - Interactively, using **pyspark**
 - Non-interactively, using **spark-submit**
 - **Within the Jupyter Notebook**
- Run interactively: open the Python version of the Spark shell
bin/pyspark
- Run non-interactively: Write Python scripts and run script:
bin/spark-submit my_script.py
- **Run in Jupyter Notebook: Assignment 5**

Installing pyspark

```
$ python
```

```
Python 3.6.5 |Anaconda, Inc.| (default, Apr 26 2018,  
08:42:37)
```

```
[GCC 4.2.1 Compatible Clang 4.0.1  
(tags/RELEASE_401/final)] on darwin
```

```
Type "help", "copyright", "credits" or "license" for  
more information.
```

```
>>> import pyspark
```

```
Traceback (most recent call last):
```



```
File "<stdin>", line 1, in <module>ModuleNotFoundError:  
No module named 'pyspark'
```

```
$ pip install pyspark
```

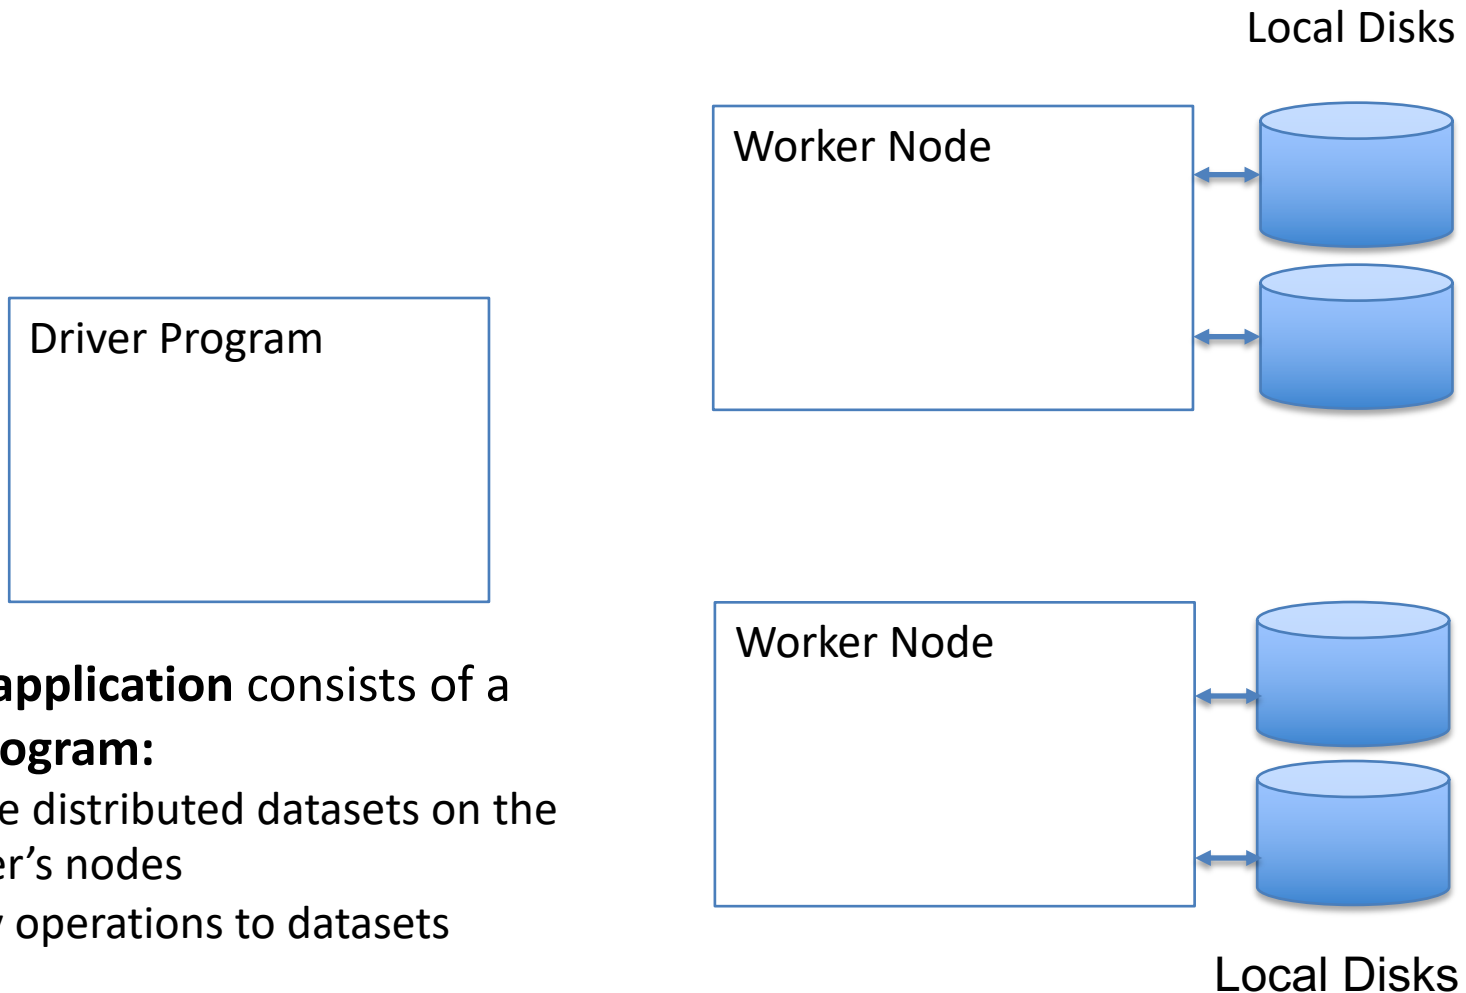
Spark Core Concepts

- A **Spark application** consists of a **driver program**
- A driver program:
 - Defines distributed datasets on the cluster
 - Applies operations to datasets
- A driver program accesses Spark through a **SparkContext** object
- A **SparkContext** represents a connection to a computing cluster
- Spark uses SparkContext to build resilient **distributed datasets (RDDs)**

Spark Core Concepts



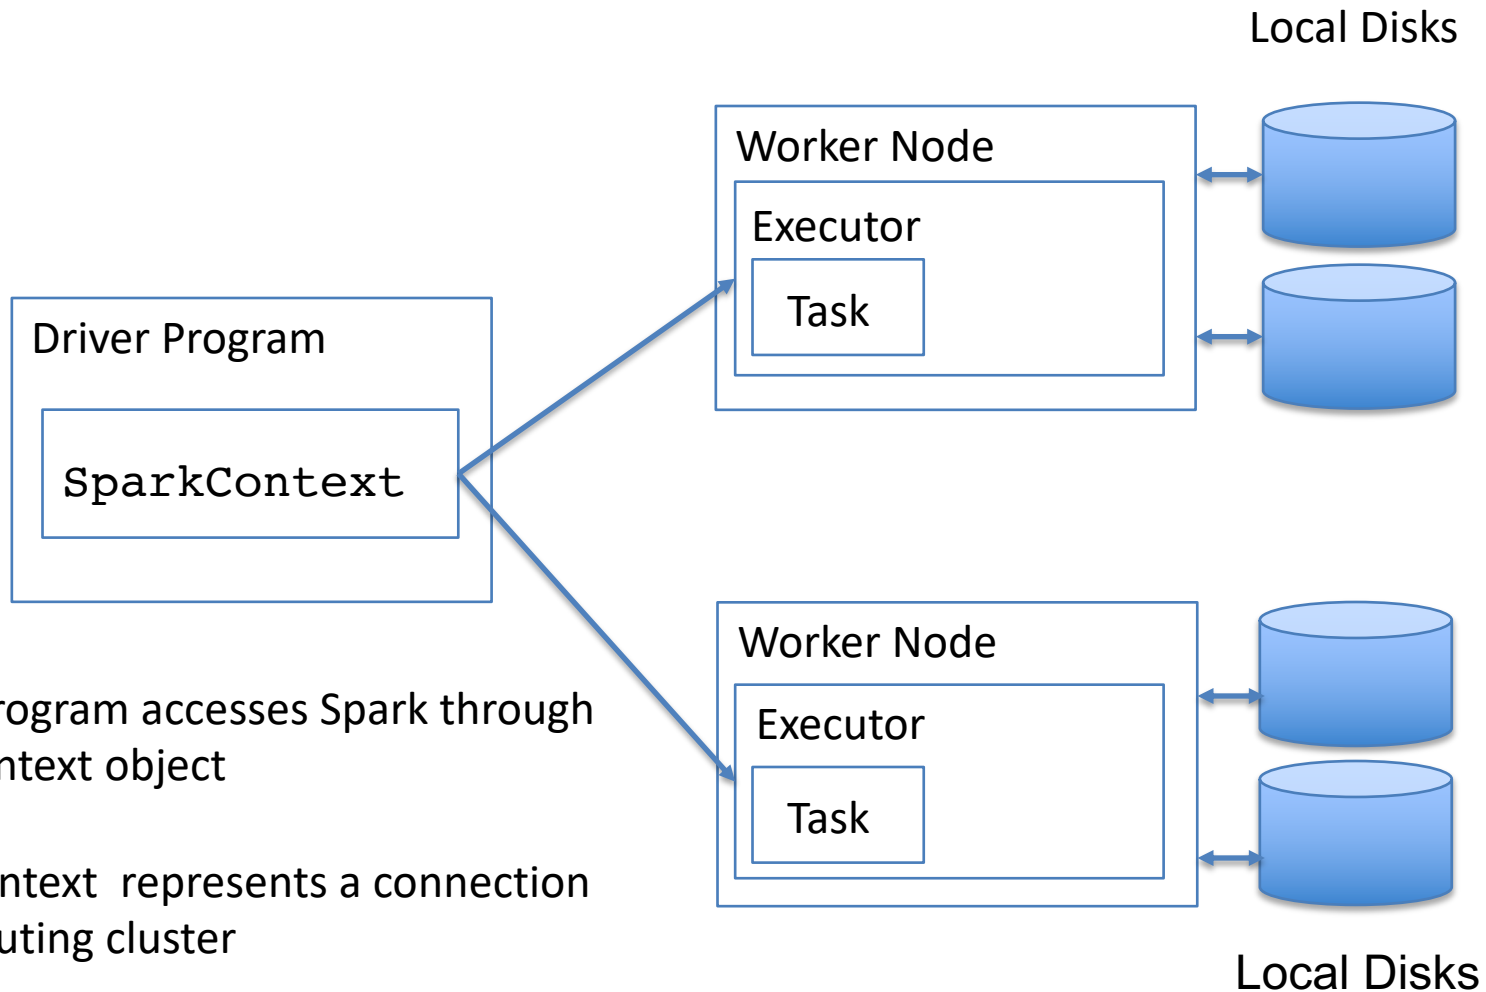
Spark Core Concepts



A **Spark application** consists of a **driver program**:

- Define distributed datasets on the cluster's nodes
- Apply operations to datasets

Spark Core Concepts



A driver program accesses Spark through a SparkContext object

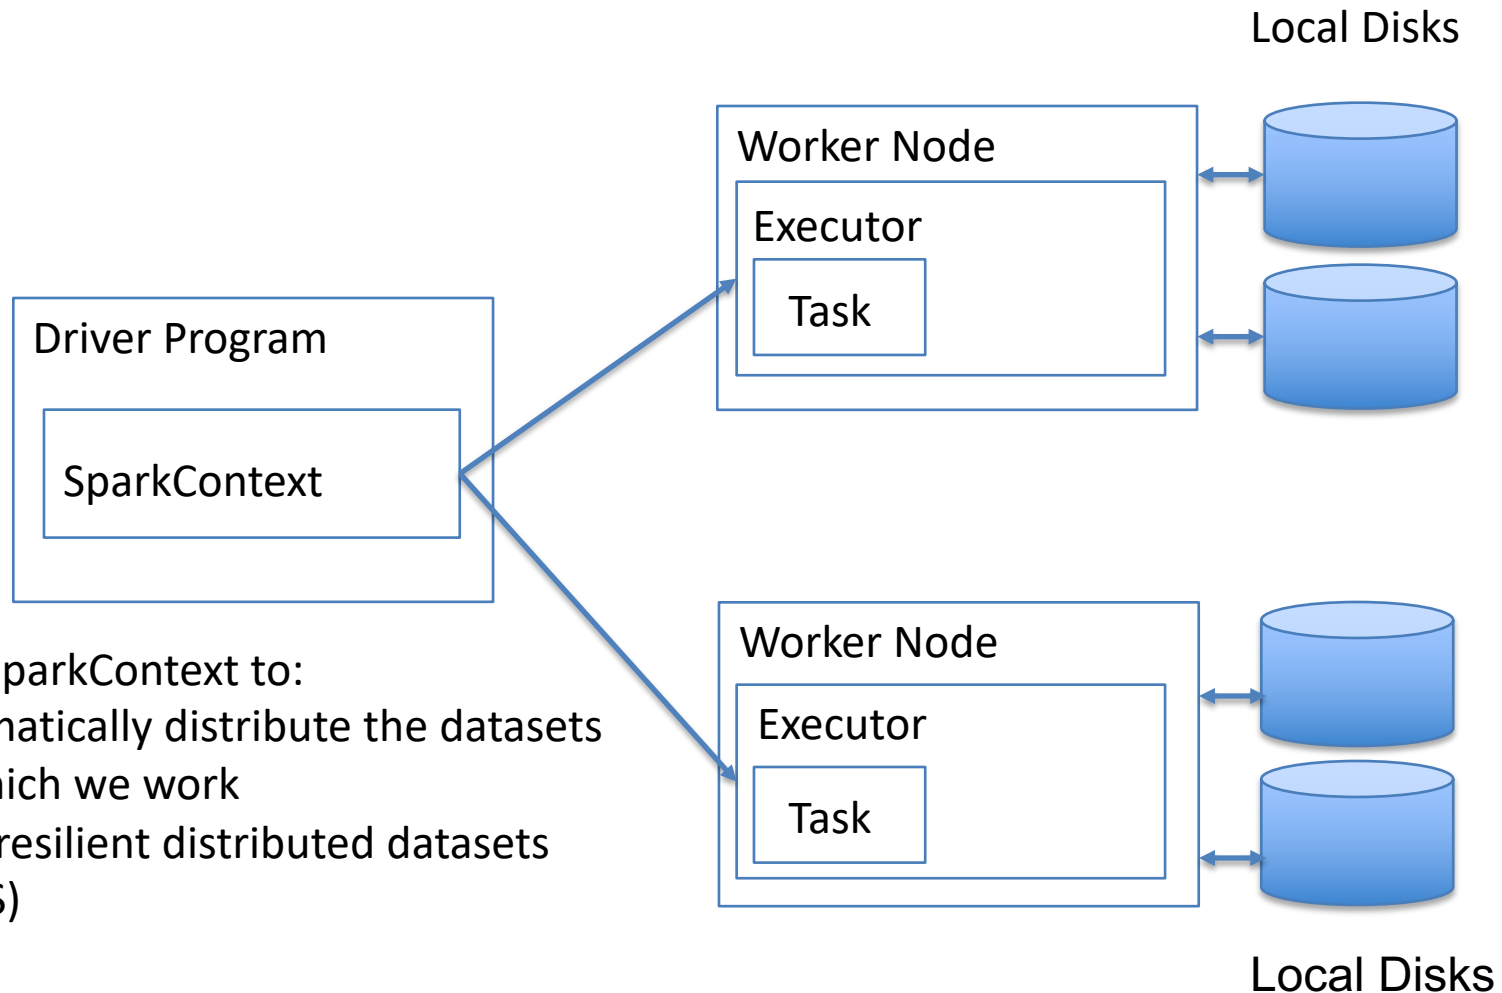
A SparkContext represents a connection to a computing cluster

Create a SparkContext

In [1]

```
from pyspark import SparkContext  
  
sc = SparkContext.getOrCreate()
```

Spark Core Concepts



We use SparkContext to:

- Automatically distribute the datasets on which we work
- Build resilient distributed datasets (RDDs)


Resilient Distributed Datasets

- Spark operates on a distributed collections of data called Resilient Distributed Datasets (or RDDs)
- We express Spark computation through operations on RDDs
- Datasets are automatically distributed across a cluster
- Operations are automatically parallelized across a cluster

→ RDDs are Spark's fundamental abstraction for distributed data and computation

Given the file "FoxInSocks.txt"

*When tweetle beetles fight,
it's called a tweetle beetle battle.
And when they battle in a puddle,
it's a tweetle beetle puddle battle.
And when tweetle beetles battle with paddles in a puddle,
They call it a tweetle beetle puddle paddle battle.*



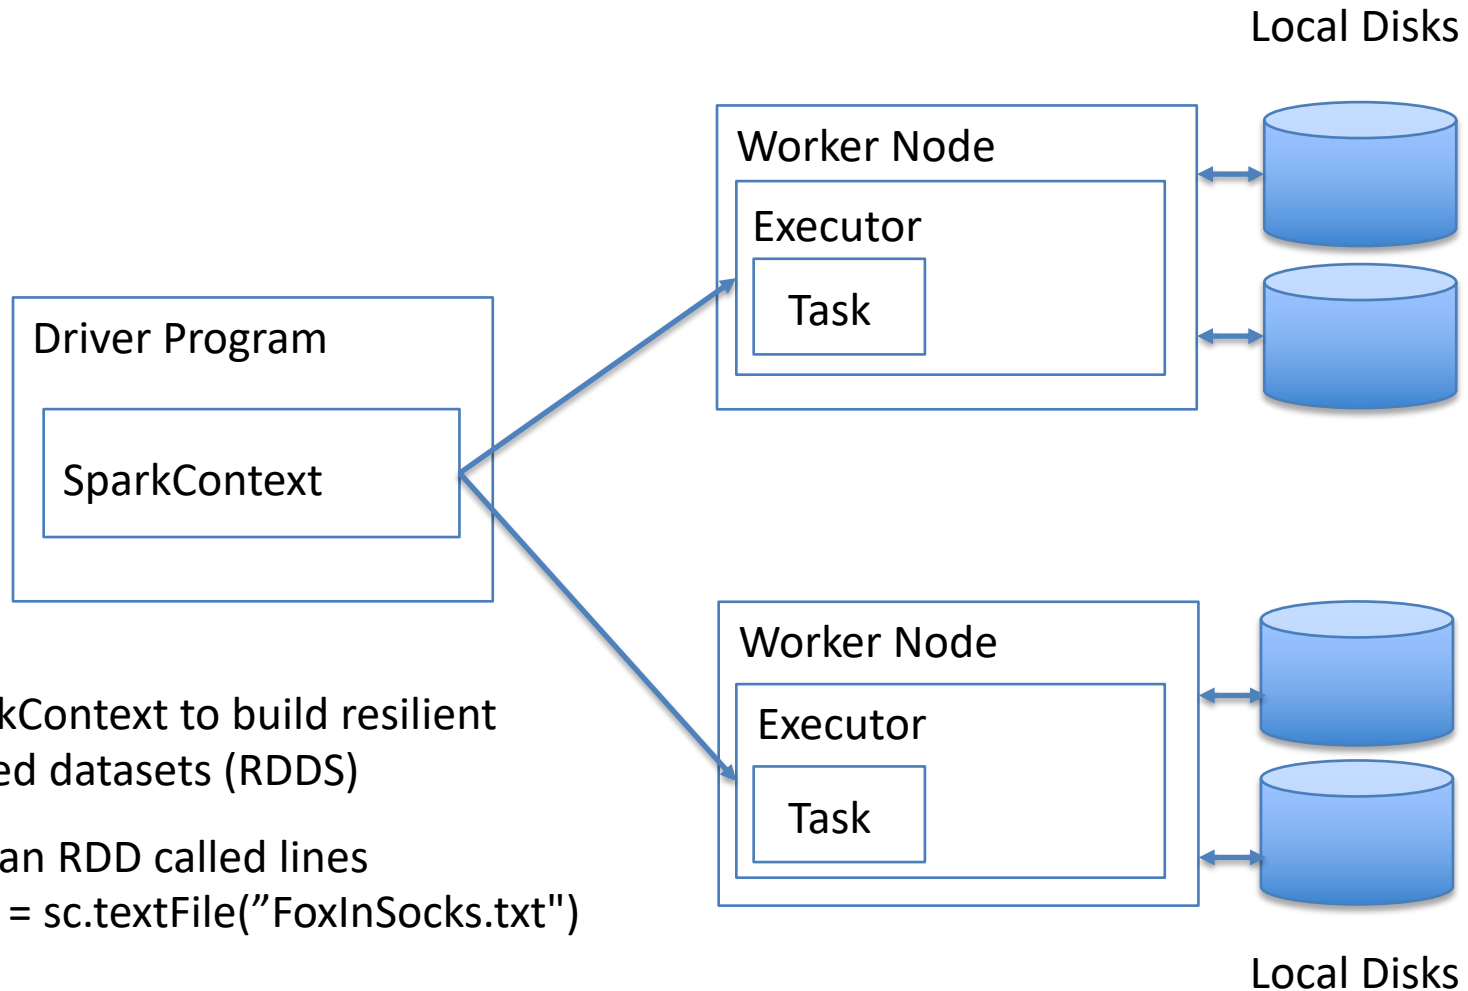
Create an RDD called lines

```
>>> lines = sc.textFile("FoxInSocks.txt")
```

File lines automatically distributed across nodes of 2-node cluster

Node 1	When tweetle beetles fight, it's called a tweetle beetle battle. And when they battle in a puddle,
Node 2	it's a tweetle beetle puddle battle. And when tweetle beetles battle with paddles in a puddle, They call it a tweetle beetle puddle paddle battle.

Spark Core Concepts



Use SparkContext to build resilient distributed datasets (RDDs)

Create an RDD called lines

```
>>> lines = sc.textFile("FoxInSocks.txt")
```


Create a SparkContext

In [1]

```
from pyspark import SparkContext
sc = SparkContext.getOrCreate()
# Load list of words
lines = sc.textFile('FoxInSocks.txt')
```

Count Number of Lines

In [1]

```
from pyspark import SparkContext
sc = SparkContext.getOrCreate()
# Load list of words
lines = sc.textFile('FoxInSocks.txt')
# Count the number of items in this RDD
print(lines.count())
```

Print First Line

In [1]

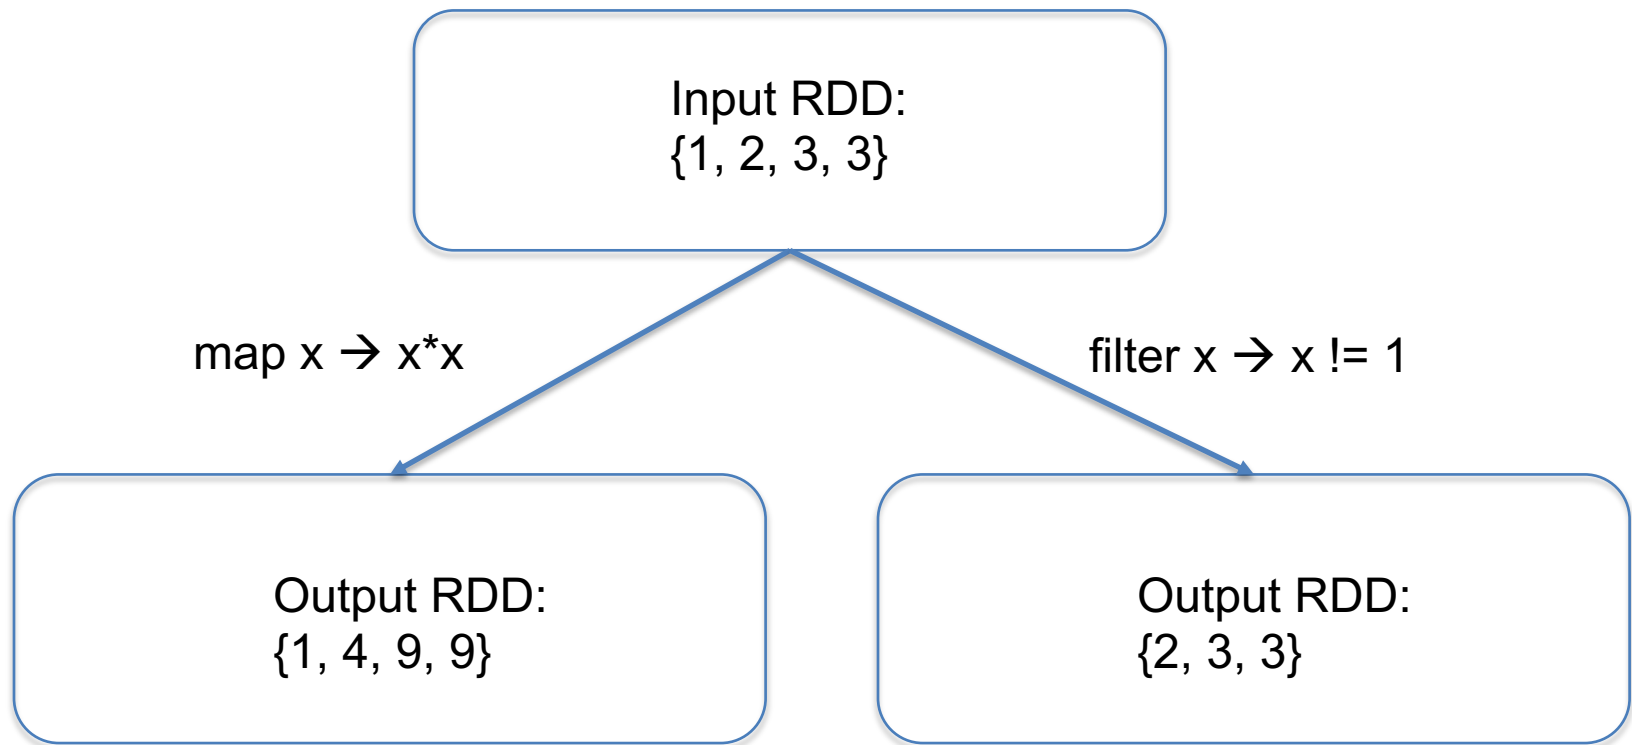
```
from pyspark import SparkContext
sc = SparkContext.getOrCreate()
# Load list of words
lines = sc.textFile('FoxInSocks.txt')
# First item in this RDD, i.e. first line of FoxInSocks.txt
print( lines.first())
```

Operations

- **Transformations:** lazily evaluated—no immediate computation
 - “Return” new RDDs obtained by transforming an old RDD
 - Input: **RDD type** → OPERATION → Output: **RDD type**
- **Actions:** cause all *queued* transformations to be applied
 - Return a list or value to the driver (serial) process
 - Input: **RDD type** → OPERATION → Output: **NOT a RDD type** (e.g., integer)

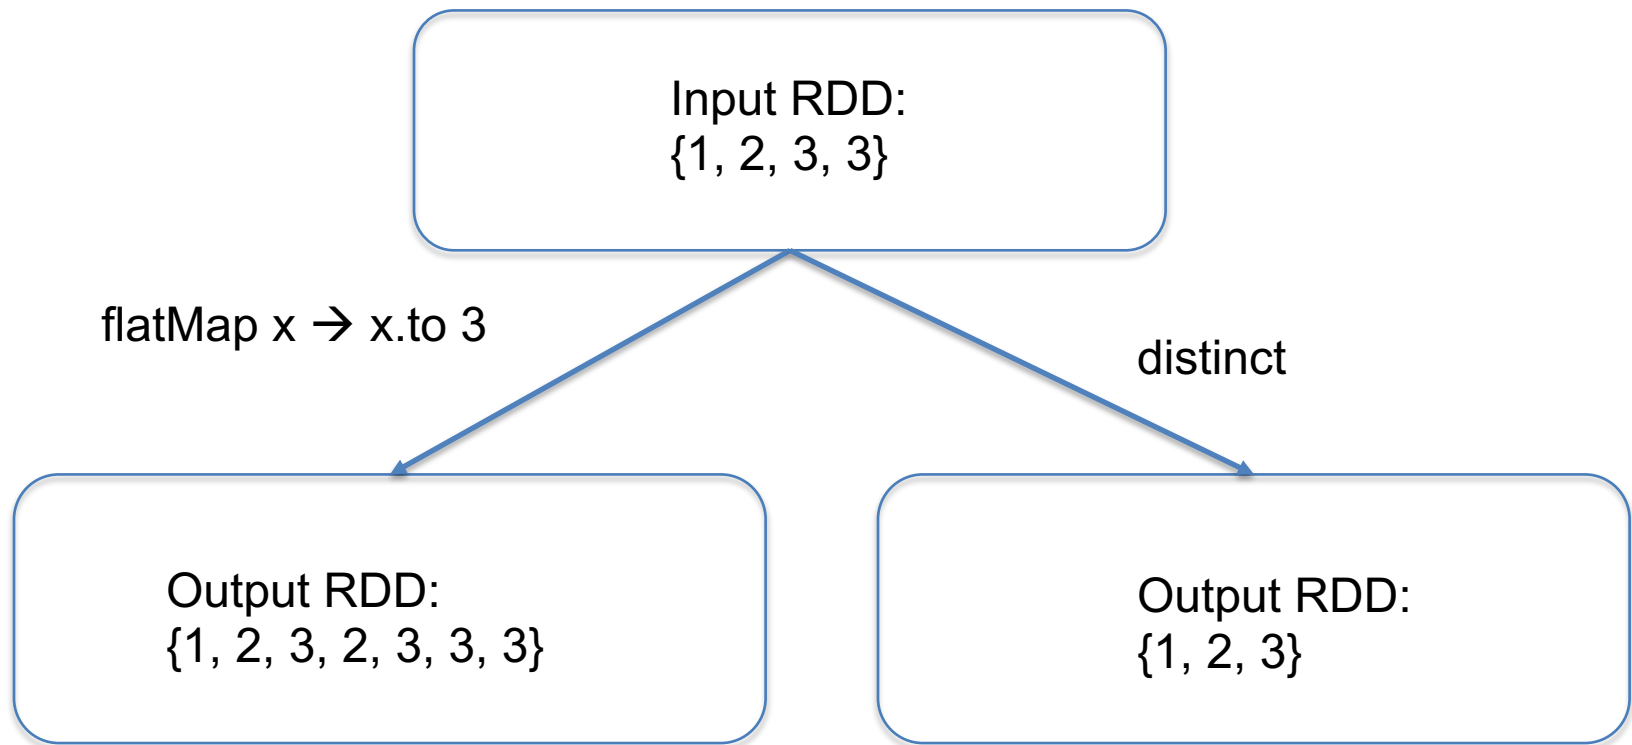
Transformations I

- Transformations (lazily evaluated—no immediate computation)



Transformations I (Cont.)

- Transformations (lazily evaluated—no immediate computation)



Transformations I (Cont.)

From Book in Chap 2

- Transformations (lazily evaluated—no immediate computation)

Function Name	Purpose	Example	Result
map	Apply a function to each element in the RDD and return an RDD of the result	<code>rdd.map(x => x + 1)</code>	<code>{2, 3, 4, 4}</code>
flatMap	Apply a function to each element in the RDD and return an RDD of the contents of the iterators returned. Often used to extract words.	<code>rdd.flatMap(x => x.to(3))</code>	<code>{1, 2, 3, 2, 3, 3, 3}</code>
filter	Return an RDD consisting of only elements which pass the condition passed to filter	<code>rdd.filter(x => x != 1)</code>	<code>{2, 3, 3}</code>
distinct	Remove duplicates	<code>rdd.distinct()</code>	<code>{1, 2, 3}</code>
<code>sample(withReplacement, fraction, [seed])</code>	Sample an RDD	<code>rdd.sample(false, 0.5)</code>	non-deterministic

Use *map* Operation on Numbers

```
In [1] from pyspark import SparkContext
      sc = SparkContext.getOrCreate()

      # create a collection
      # elements form distributed dataset that can be operated on in parallel
      # Spark set number of partitions automatically based on cluster
      # you can set partitions manually by passing number
      # second parameter in parallelize (e.g., sc.parallelize(data, 10))
      numbers = sc.parallelize([1, 2, 3, 3])
      squared = numbers.map(lambda x: x * x)
```


Use *map* and *first* Operations on Text

In [1]

```
from pyspark import SparkContext  
sc = SparkContext.getOrCreate()  
  
lines = sc.parallelize(["hello world", "hi"])  
words = lines.map(lambda line: line.split(" "))  
words.first()
```

What is the output?

Use *map* and *first* Operations on Text

In [1]

```
from pyspark import SparkContext  
sc = SparkContext.getOrCreate()  
  
lines = sc.parallelize(["hello world", "hi"])  
words = lines.map(lambda line: line.split(" "))  
words.first()
```

```
[['hello', 'world'], ['hi']]
```

Use *flatMap* and *first* Operations on Text

In [1]

```
from pyspark import SparkContext  
sc = SparkContext.getOrCreate()  
  
lines = sc.parallelize(["hello world", "hi"])  
words = lines.flatMap(lambda line: line.split(" "))  
words.first()
```

What is the output?

Use *flatMap* and *first* Operations on Text

In [1]

```
from pyspark import SparkContext
sc = SparkContext.getOrCreate()

lines = sc.parallelize(["hello world", "hi"])
words = lines.flatMap(lambda line: line.split(" "))
words.first()
```

```
['hello', 'world', 'hi']
```

Work on Text with *filter*

In [2]

```
from pyspark import SparkContext
sc = SparkContext.getOrCreate()
lines = sc.textFile('FoxInSocks.txt')

def hasWhen(line):
    return 'when' in line

whenLines = lines.filter(hasWhen)
```

Work on Text with *filter* (Cont.)

In [3]

```
from pyspark import SparkContext  
sc = SparkContext.getOrCreate()  
lines = sc.textFile('FoxInSocks.txt')  
  
whenLines = lines.filter(lambda line: 'when' in line)
```

Transformations II (Cont.)

From Book in Chap 2

- Transformations (lazily evaluated—no immediate computation)

RDD1
{coffee, coffee, panda ,
monkey, tea }

RDD2
{coffee, monkey, kitty}

RDD1.distinct()
{coffee, panda,
monkey, tea}

RDD1.union(RDD2)
{coffee, coffee, coffee,
panda, monkey,
monkey, tea, kitty}

RDD1.intersection(RDD2)
{coffee, monkey}

RDD1.subtract(RDD2)
{panda, tea}

Transformations II (Cont.)

From Book in Chap 2

- Transformations (lazily evaluated—no immediate computation)

RDDs for the examples in the table:

`rdd = {1, 2, 3}`

`other = {3, 4, 5}`

Function Name	Purpose	Example	Result
<code>union</code>	Produce an RDD contain elements from both RDDs	<code>rdd.union(other)</code>	<code>{1, 2, 3, 3, 4, 5}</code>
<code>intersection</code>	RDD containing only elements found in both RDDs	<code>rdd.intersection(other)</code>	<code>{3}</code>
<code>subtract</code>	Remove the contents of one RDD (e.g. remove training data)	<code>rdd.subtract(other)</code>	<code>{1, 2}</code>
<code>cartesian</code>	Cartesian product with the other RDD	<code>rdd.cartesian(other)</code>	<code>{(1, 3), (1, 4), ... (3,5)}</code>

Actions

From Book in Chap 2

RDD for the examples in the table:

`rdd = {1, 2, 3, 3}`

Function Name	Purpose	Example (In Scala)	Result
<code>collect()</code>	Return all elements from the RDD	<code>rdd.collect()</code>	<code>{1, 2, 3, 3}</code>
<code>count()</code>	Number of elements in the RDD	<code>rdd.count()</code>	<code>4</code>
<code>take(num)</code>	Return num elements from the RDD	<code>rdd.take(2)</code>	<code>{1, 2}</code>
<code>top(num)</code>	Return the top num elements the RDD	<code>rdd.top(2)</code>	<code>{3, 3}</code>
<code>takeOrdered(num)(ordering)</code>	Return num elements based on providing ordering	<code>rdd.takeOrdered(2)(myOrdering)</code>	<code>{3, 3}</code>

Function Name	Purpose	Example (In Scala)	Result
takeSample(withReplacement, num, [seed])	Return num elements at random	<code>rdd.takeSample(false, 1)</code>	non-deterministic
reduce(func)	Combine the elements of the RDD together in parallel (e.g. sum)	<code>rdd.reduce((x, y) => x + y)</code>	9
fold(zero)(func)	Same as reduce but with the provided zero value	<code>rdd.fold(0)((x, y) => x + y)</code>	9
aggregate(zeroValue)(seqOp, combOp)	Similar to reduce but used to return a different type	<code>rdd.aggregate(0, 0)({case (x, y) => (y._1() + x, y._2() + 1)}, {case (x, y) => (y._1() + x._1(), y._2() + x._2())})</code>	(9, 4)
foreach(func)	Apply the provided function to each element of the RDD	<code>rdd.foreach(func)</code>	nothing

REVIEW:

Use *map* Operation on Numbers

```
In [1] from pyspark import SparkContext
sc = SparkContext.getOrCreate()

# create a collection
# elements form distributed dataset that can be operated on in parallel
# Spark set number of partitions automatically based on cluster
# you can set partitions manually by passing number
# second parameter in parallelize (e.g., sc.parallelize(data, 10))
numbers = sc.parallelize([1, 2, 3, 3])
squared = numbers.map(lambda x: x * x)
squared.collect()
```

```
[1, 4, 9, 9]
```

REVIEW:

Use *flatMap* Operation on Numbers

```
In [1] from pyspark import SparkContext
       sc = SparkContext.getOrCreate()

       lines = sc.parallelize(["hello world", "hi"])
       words = lines.flatMap(lambda line: line.split(" "))
       words.collect()
```

```
['hello', 'world', 'hi']
```

REVIEW:

Work on Text with *filter*

In [3]

```
from pyspark import SparkContext  
sc = SparkContext.getOrCreate()  
lines = sc.textFile('FoxInSocks.txt')  
  
whenLines = lines.filter(lambda line: 'when' in line)  
whenLines.collect()
```

```
['And when they battle in a puddle, ', 'AND when tweetle beetles  
battle with paddles in a puddle, ']
```

Use *map* and *collect* Operations on Numbers

In [1]

```
from pyspark import SparkContext
sc = SparkContext.getOrCreate()

numbers = sc.parallelize([1, 2, 3, 4])
squared = numbers.map(lambda x: x * x).collect()
for num in squared:
    print (num)
```

What is the output?

Create Key-Values in RDDs (I)

In [3]

```
from pyspark import SparkContext  
sc = SparkContext.getOrCreate()  
lines = sc.textFile('FoxInSocks.txt')  
  
pairs= lines.map(lambda x: (x.split(" ")[0], x))
```

Create Key-Values in RDDs (I)

In [3]

```
from pyspark import SparkContext
sc = SparkContext.getOrCreate()
lines = sc.textFile('FoxInSocks.txt')

pairs= lines.map(lambda x: (x.split(" ")[0], x))

<When, When tweetle beetles fight,>
<it's, it's called a tweetle beetle battle.>
<And, And when they battle in a puddle,>
....
```


Create Key-Values in RDDs (II)

In [3]

```
from pyspark import SparkContext  
sc = SparkContext.getOrCreate()  
lines = sc.textFile("FoxInSocks.txt")  
  
pairs= lines.map(lambda x: (x.split(" ")[0], x))  
results = pairs.filter(lambda x: len(x[1]) < 28)
```

Create Key-Values in RDDs

In [3]

```
from pyspark import SparkContext
```

```
sc = SparkContext.getOrCreate()
```

```
lines = sc.textFile("FoxInSocks.txt")
```

```
pairs= lines.map(lambda x: (x.split(" ")[0], x))
```

```
results = pairs.filter(lambda x: len(x[1]) < 28)
```

<When, When tweetle beetles fight,>

...

Create Key-Values in RDDs

In [3]

```
from pyspark import SparkContext  
sc = SparkContext.getOrCreate()  
lines = sc.textFile('FoxInSocks.txt')  
  
words = lines.flatMap(lambda x: x.split(" "))
```

Create Key-Values in RDDs

In [3]

```
from pyspark import SparkContext  
sc = SparkContext.getOrCreate()  
lines = sc.textFile('FoxInSocks.txt')  
  
words = lines.flatMap(lambda x: x.split(" "))  
pairs= words.map(lambda x: (x, 1))
```

Create Key-Values in RDDs

In [3]

```
from pyspark import SparkContext
```

```
sc = SparkContext.getOrCreate()
```

```
lines = sc.textFile('FoxInSocks.txt')
```

```
words = lines.flatMap(lambda x: x.split(" "))
```

```
pairs = words.map(lambda x: (x, 1))
```

```
<"When", 1> <"tweetle", 1> <"beetles", 1> <"fight", 1>  
...
```

Create Key-Values in RDDs

In [3]

```
from pyspark import SparkContext
sc = SparkContext.getOrCreate()
lines = sc.textFile("FoxInSocks.txt")

words = lines.flatMap(lambda x: x.split(" "))
pairs = words.map(lambda x: (x, 1))
results = pairs.reduceByKey(lambda x, y: x + y)
```

Create Key-Values in RDDs

In [3]

```
from pyspark import SparkContext
sc = SparkContext.getOrCreate()
lines = sc.textFile("FoxInSocks.txt")

words = lines.flatMap(lambda x: x.split(" "))
pairs = words.map(lambda x: (x, 1))
results = pairs.reduceByKey(lambda x, y: x + y)

## TRY to print the results
print(results)
```

Create Key-Values in RDDs

In [3]

```
from pyspark import SparkContext
sc = SparkContext.getOrCreate()
lines = sc.textFile("FoxInSocks.txt")

words = lines.flatMap(lambda x: x.split(" "))
pairs = words.map(lambda x: (x, 1))
results = pairs.reduceByKey(lambda x, y: x + y)

## TRY to print the results
print(results)
ERRORS!!!
```


The Special Case of *ReduceByKey*

- Reduce takes a function and use it to combine values
- ReduceByKey takes a function and use it to combine values
- **BUT** ReduceByKey **DO NOT** implemented as an action
 - Return a new RDD consisting of each key and the reduced value for that key

WHY?

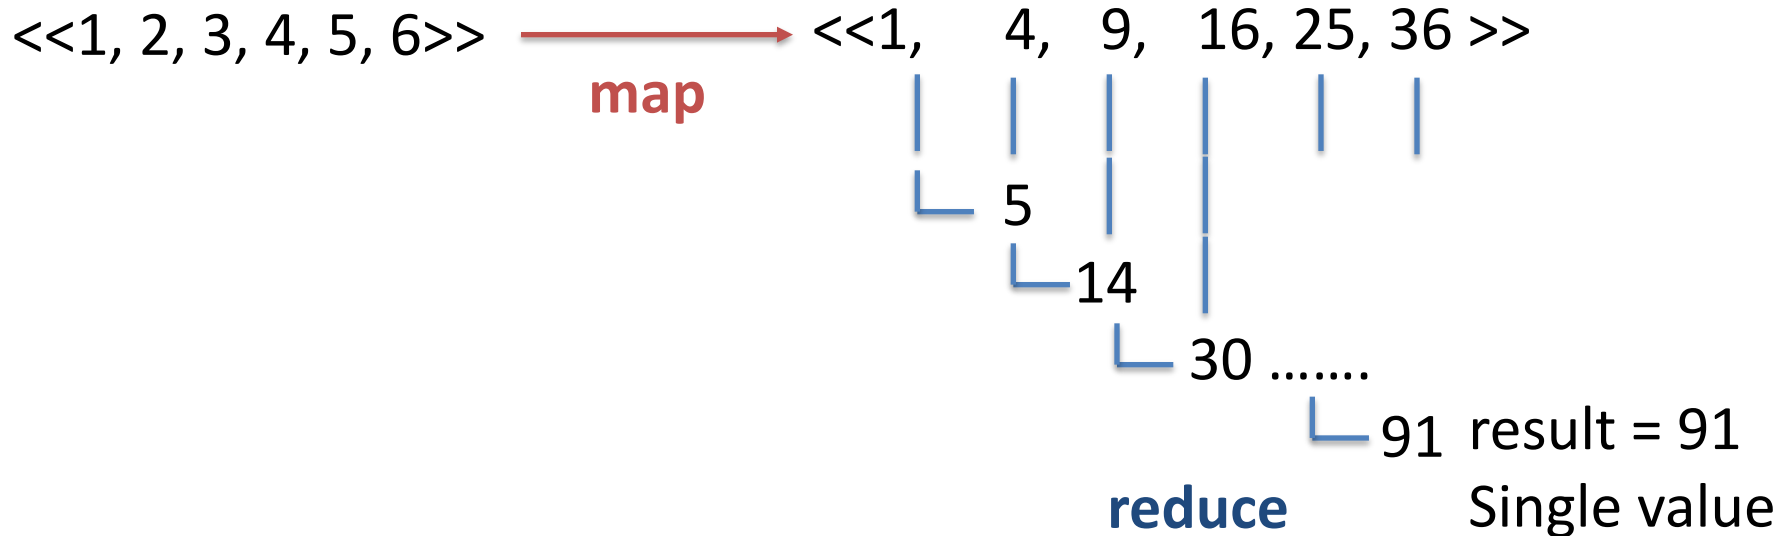
The Special Case of *ReduceByKey*

- Reduce takes a function and use it to combine values
- ReduceByKey takes a function and use it to combine values
- **BUT** ReduceByKey **DO NOT** implemented as an action
 - Return a new RDD consisting of each key and the reduced value for that key

WHY?

- *reduceByKey* runs several parallel reduce operations:
 - one for each key in the dataset
 - each operation combines values together which have the same key.
- **Datasets can have very large numbers of keys!!!!**

Use *reduce* Operation on Numbers (I)




Note: Spark does **NOT** guarantee the order of operands

Use *reduce* Operation on Numbers I

In [1]

```
from pyspark import SparkContext
sc = SparkContext.getOrCreate()

numbers = sc.parallelize([1, 2, 3, 4, 5, 6])
squared = numbers.map(lambda x: x * x)
result = squared.reduce(lambda x, y: x+y) # single value
```



Use *reduce* Operation on Numbers I

This is an RDD

<< (B, 5), (B, 4),
 (A,2), (B,3),
 (A,1), (C, 0) >>

→

<< (B, 5), (B, 4),
 (A,2), (B,3),
 (A,1), (C, 0) >>

results = lists.**reduceByKey**(lambda x, y: x+y)

<< (B, 12),
 (A,3),
 (C, 0) >>

←

B [5, 4, 3] = 12
A [2, 1] = 3
C [0] = 0

This is an RDD

Assignment 4

Assignment 4 - CS 594 / CS 690

- Python has map and reduce functions:
 - Do NOT take advantage of parallel processing (i.e., they are sequential)
 - Define three methods mapSequential
 - reduceSequential
 - reduceByKeySequential
- Extend Python's map and reduce functions to act like those in *Apache Spark*

Deadline: October 1 - 8AM ET

Assignment 4 - CS 690

- Read paper “*Spark: Cluster Computing with Working Sets*” Matei Zaharia, Mosharaf Chowdhury, Michael J. Franklin, Scott Shenker, Ion Stoica University of California, Berkeley
- Submit summary:
 - Add summary to your private GitHub repository
 - Use the template provided
 - Follow mandatory requirements for your summary

Deadline: October 8 - 8AM ET



THE UNIVERSITY OF
TENNESSEE
KNOXVILLE

BIG ORANGE. BIG IDEAS.®